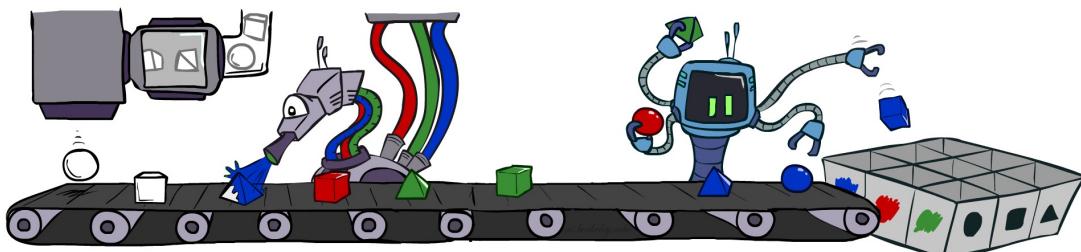


# Bayes Net Sampling

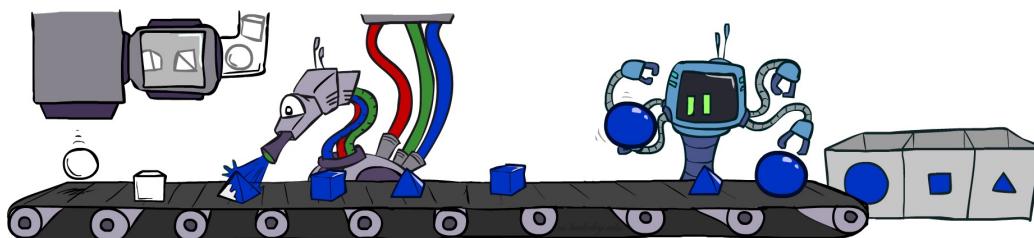
- Prior Sampling  $P$  :

- Generate complete samples from  $P(x_1, \dots, x_n)$



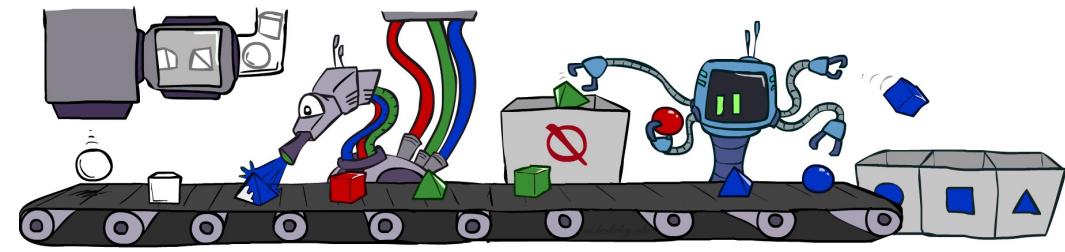
- Likelihood Weighting  $P(Q | e)$  :

- Weight samples by how well they predict  $e$



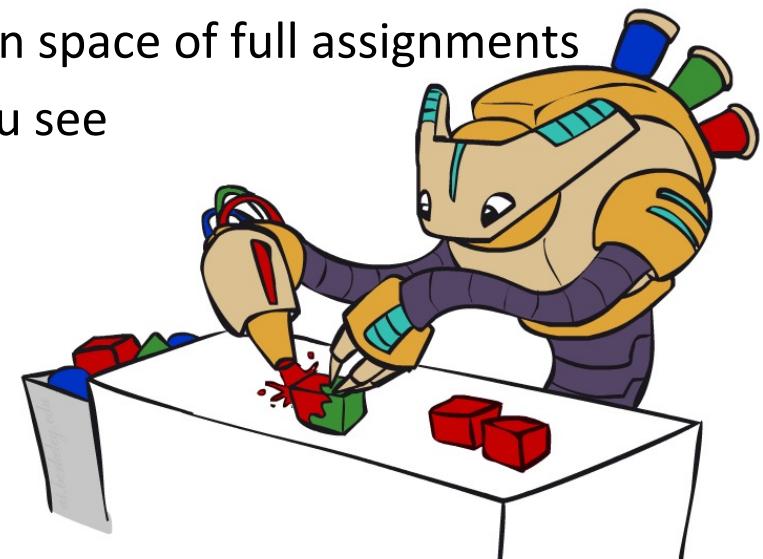
- Rejection Sampling  $P(Q | e)$  :

- Reject samples that don't match  $e$



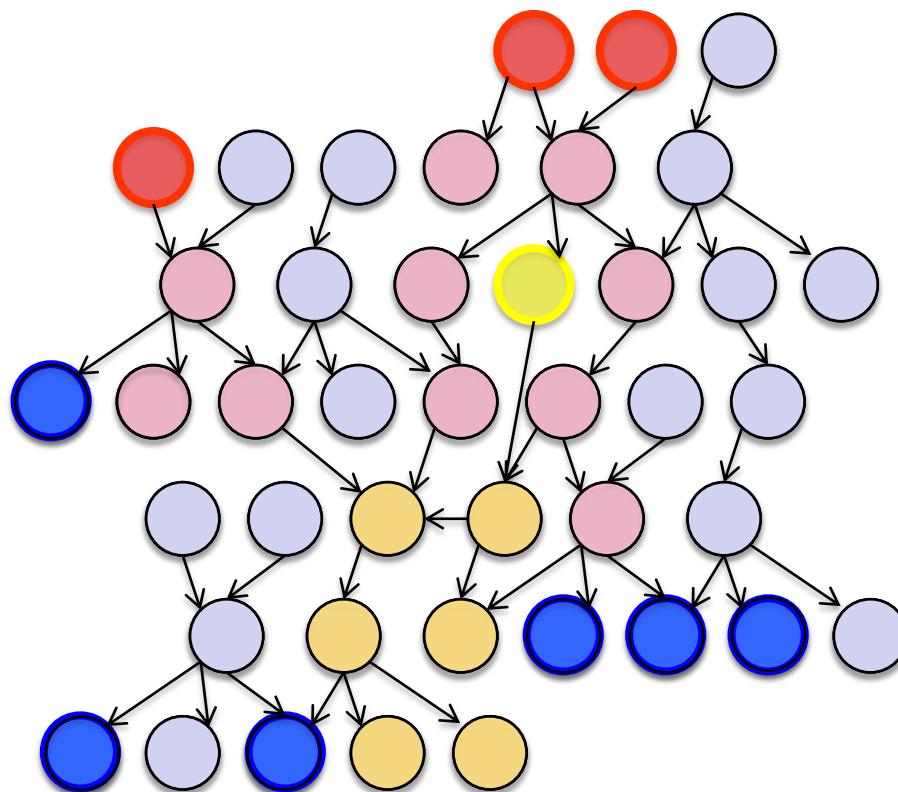
- Gibbs sampling  $P(Q | e)$  :

- Wander around in space of full assignments
- Average what you see



# Likelihood weighting

- Likelihood weighting is good
  - All samples are used
  - The values of *downstream* variables are influenced by *upstream* evidence

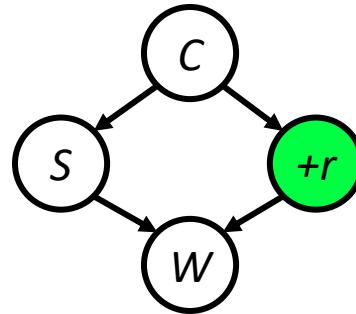


- Likelihood weighting still has weaknesses
  - The values of *upstream* variables are unaffected by *downstream* evidence
  - With evidence in  $k$  leaf nodes, weights will be  $O(2^{-k})$
  - With high probability, one lucky sample will have much larger weight than the others, dominating the result
- We would like each variable to “see” *all* the evidence!

# Gibbs Sampling Example: $P(S | +r)$

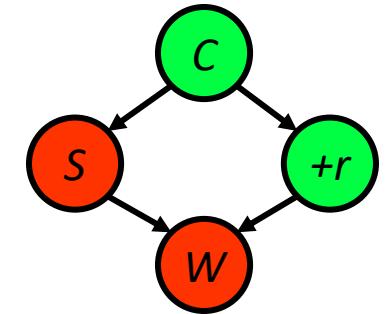
- Step 1: Fix evidence

- $R = +r$



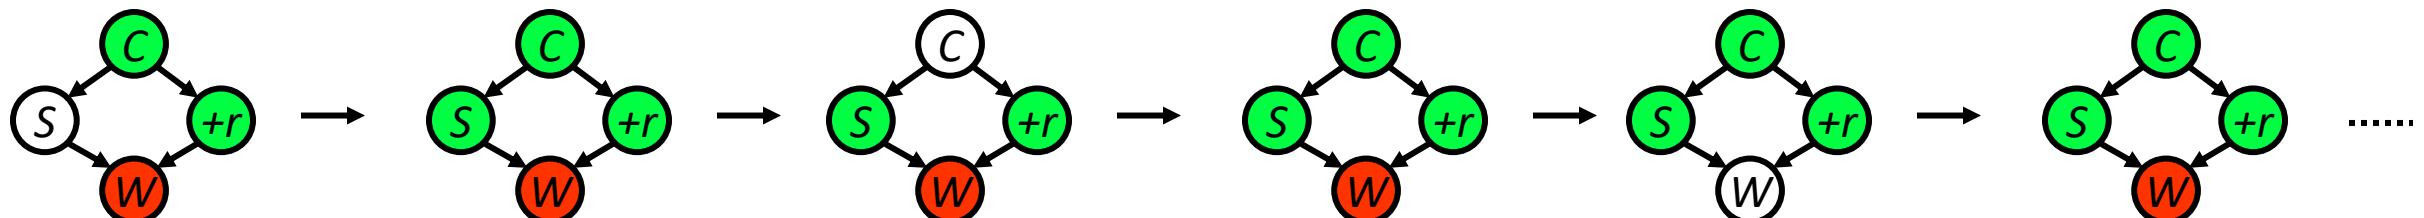
- Step 2: Initialize other variables

- Randomly



- Steps 3: Repeat

- Choose a *non-evidence* variable  $X$
  - *Resample  $X$  from  $P(X | \text{all other variables})^*$*



Sample from  $P(S | +c, -w, +r)$

Sample from  $P(C | +s, -w, +r)$

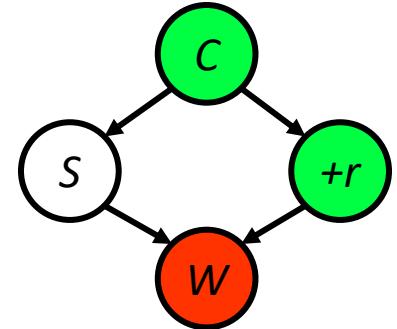
Sample from  $P(W | +s, +c, +r)$

*How to compute this?*

# Resampling Probability

- Sample from  $P(S | +c, +r, -w)$

$$\begin{aligned} P(S | +c, +r, -w) &= \frac{P(S, +c, +r, -w)}{P(+c, +r, -w)} && \text{Normalization trick} \\ &= \frac{P(S, +c, +r, -w)}{\sum_s P(s, +c, +r, -w)} && \text{Definition of marginal} \\ &= \frac{P(+c)P(S | +c)P(+r | +c)P(-w | S, +r)}{\sum_s P(+c)P(s | +c)P(+r | +c)P(-w | s, +r)} && \text{BN joint probabilities} \\ &= \frac{P(+c)P(S | +c)P(+r | +c)P(-w | S, +r)}{P(+c)P(+r | +c) \sum_s P(s | +c)P(-w | s, +r)} \\ &= \frac{P(S | +c)P(-w | S, +r)}{\sum_s P(s | +c)P(-w | s, +r)} && \text{Sum early} \quad \text{Only factors with } S \text{ remain!} \end{aligned}$$



- More generally: join all factors involving *the resampled variable* and normalize
- Better way:  $P(X_i | x_1, \dots, x_{i-1}, x_{i+1}, \dots, x_n) = P(X_i | \text{markov\_blanket}(X_i))$

# Gibbs sampling algorithm

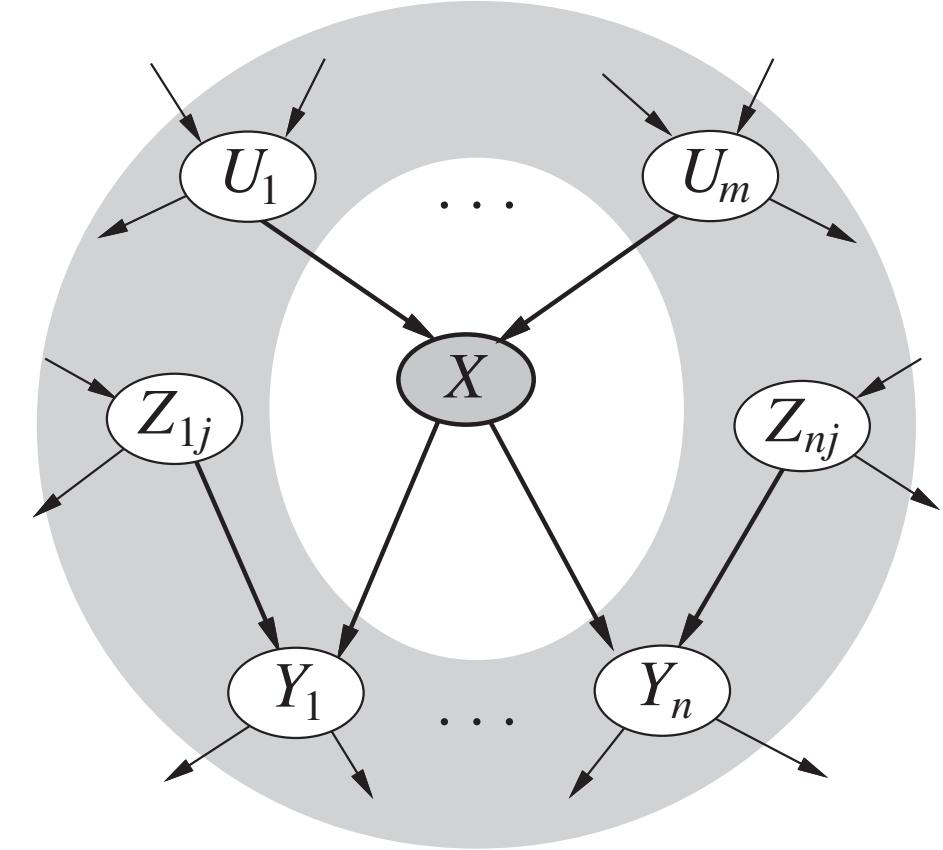
- Repeat many times

- Sample a non-evidence variable  $X_i$  from

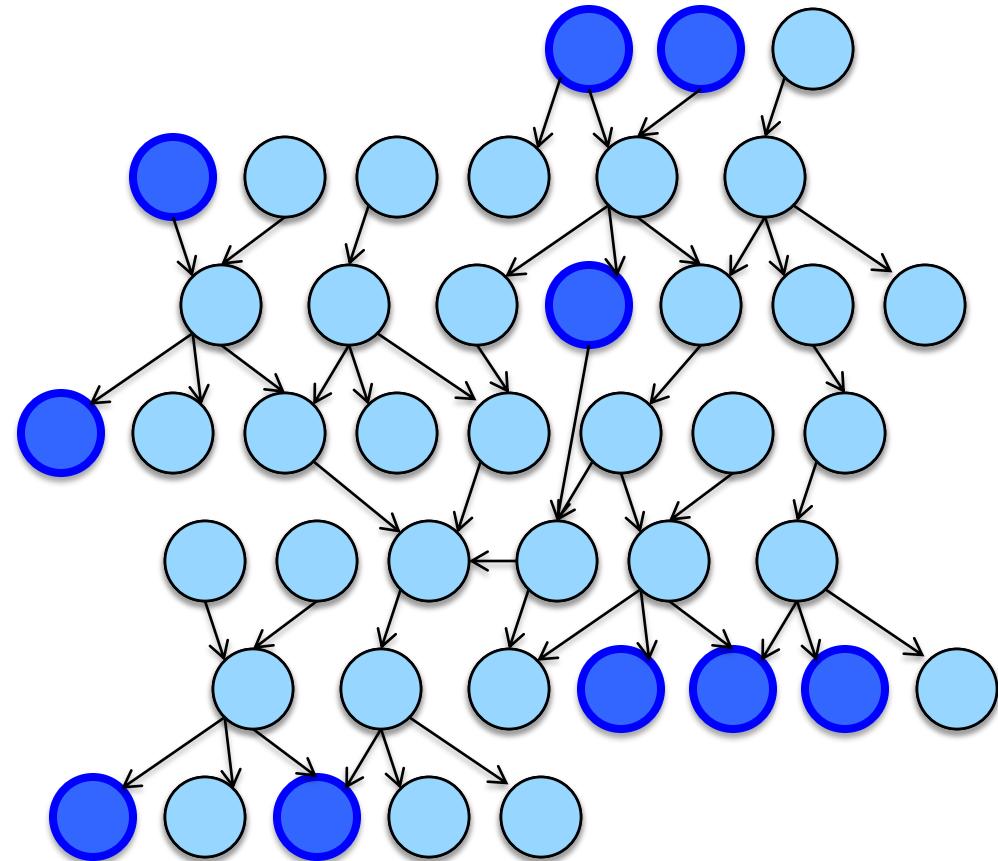
$$\begin{aligned} P(X_i | x_1, \dots, x_{i-1}, x_{i+1}, \dots, x_n) &= P(X_i | \text{markov\_blanket}(X_i)) \\ &= \alpha P(X_i | \text{parents}(X_i)) \prod_j P(y_j | \text{parents}(Y_j)) \end{aligned}$$

*join all factors involving  $X_i$  and normalize*

- Influence of evidences propagates locally, and gradually becomes global



# Advantages of Gibbs Sampling



- Samples soon begin to reflect all the evidence in the network
- Eventually samples will be drawn from the true posterior, under moderate conditions:
  - Intuitively, all variables should eventually be resampled for infinite number of times
  - -> Gibbs sampling is *consistent*!

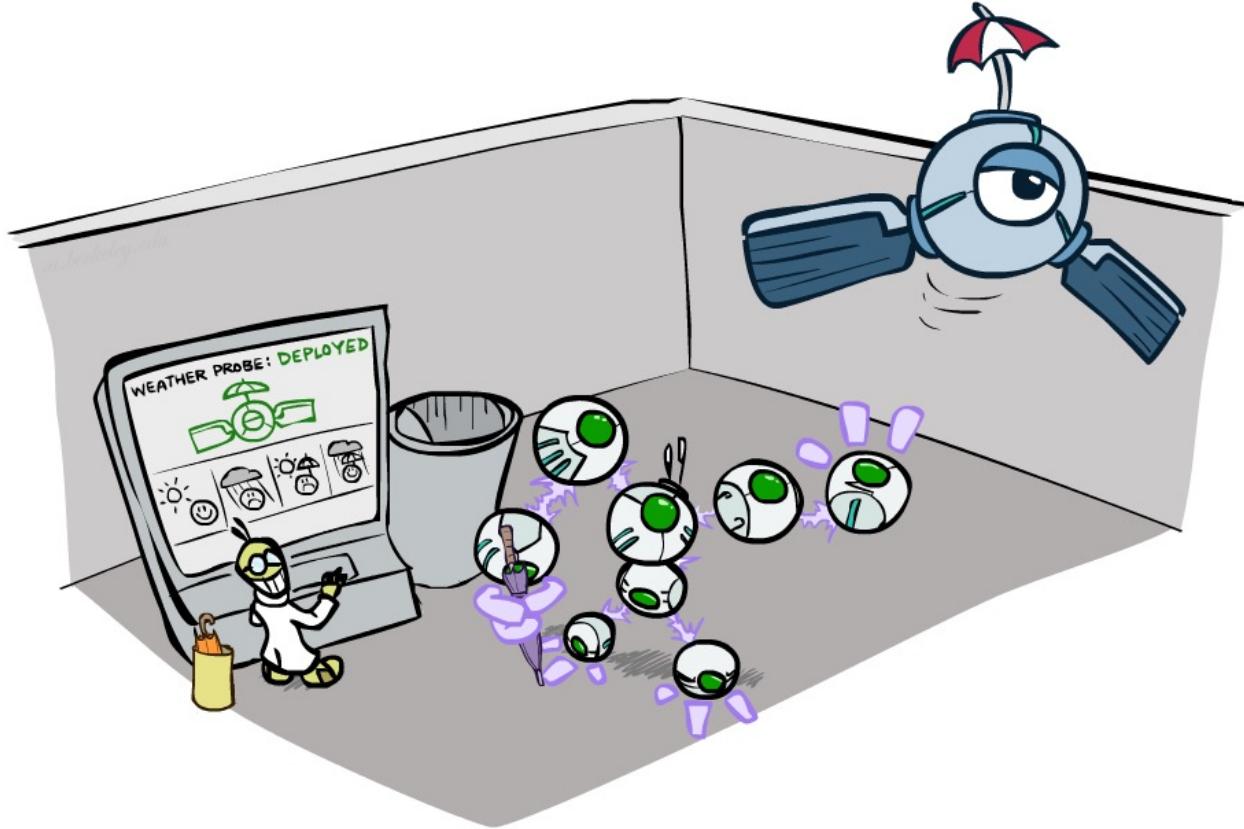
# Gibbs sampling in practice

---

- Gibbs sampling belongs to a family of sampling methods called Markov chain Monte Carlo (MCMC)
  - Specifically, it is a special case of a subset of MCMC methods called Metropolis-Hastings
- The most commonly used method for large Bayes nets
- Can be compiled to run very fast
  - Eliminate all data structure references, just multiply and sample
  - ~100 million samples per second on a laptop
- Parallelization (one processor per variable)
- Many cognitive scientists suggest the brain runs on MCMC

# CS 3317: Artificial Intelligence

## Decision Networks and Value of Information

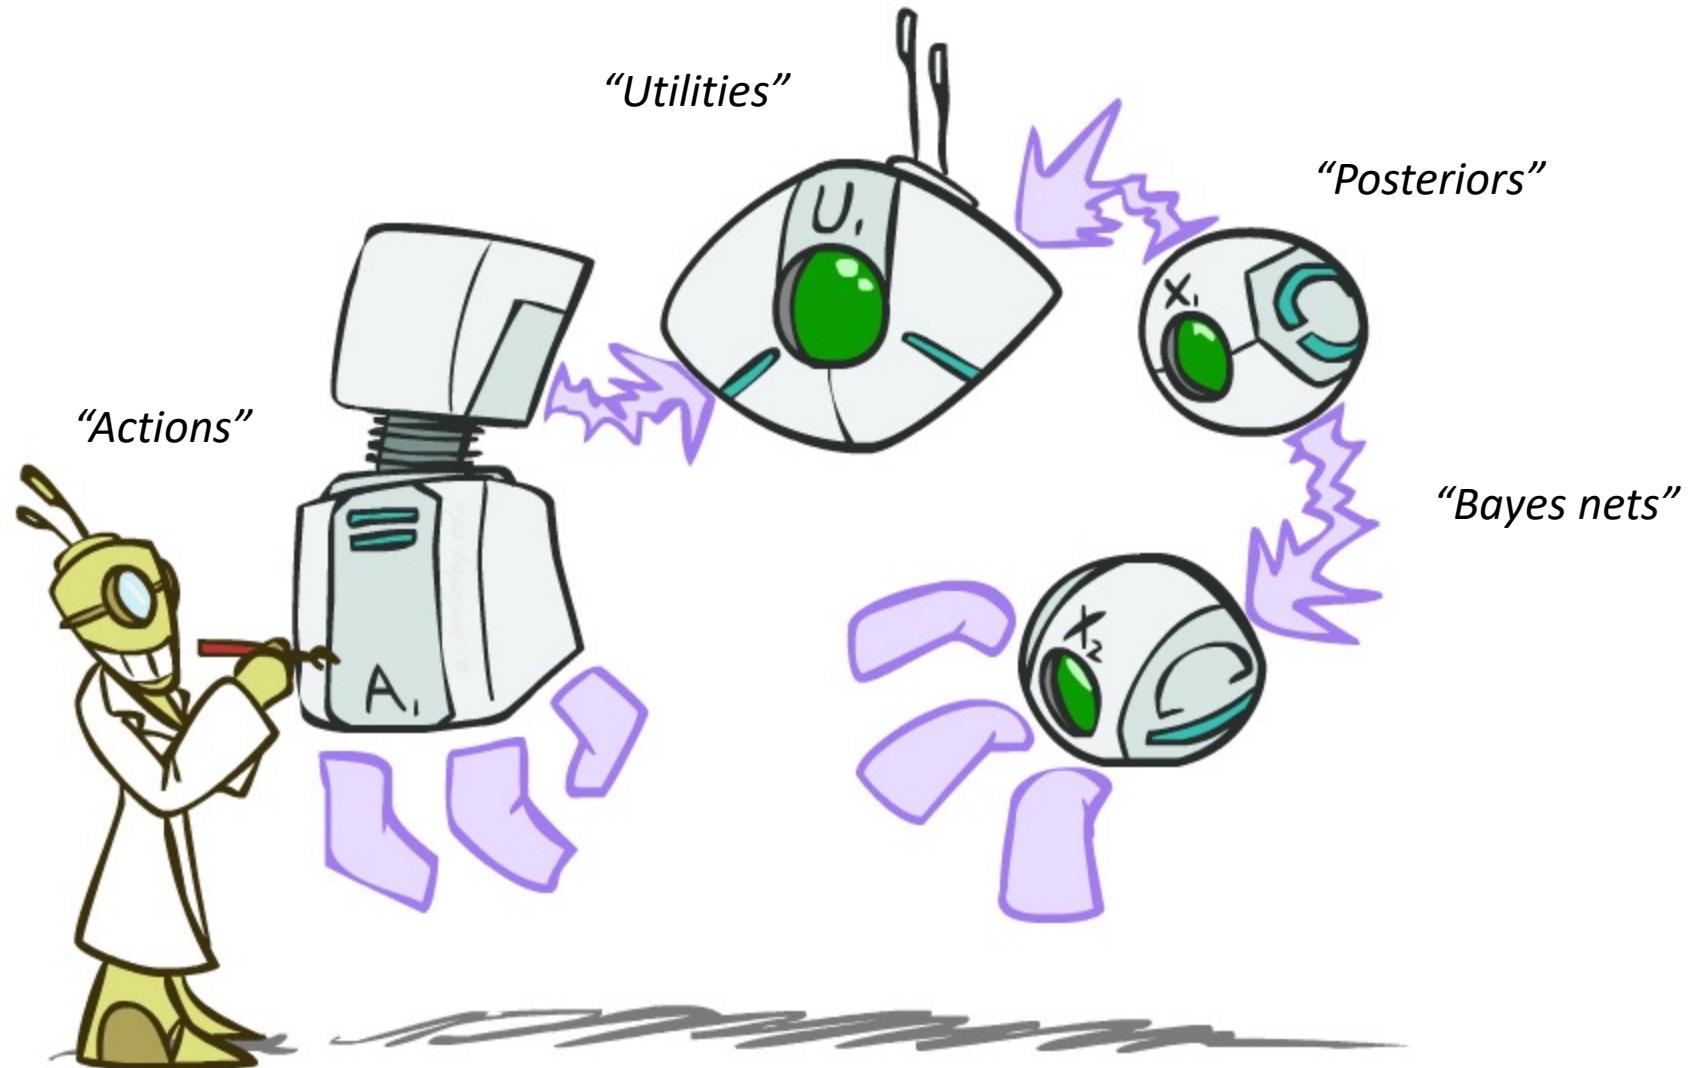


Instructor: Panpan Cai

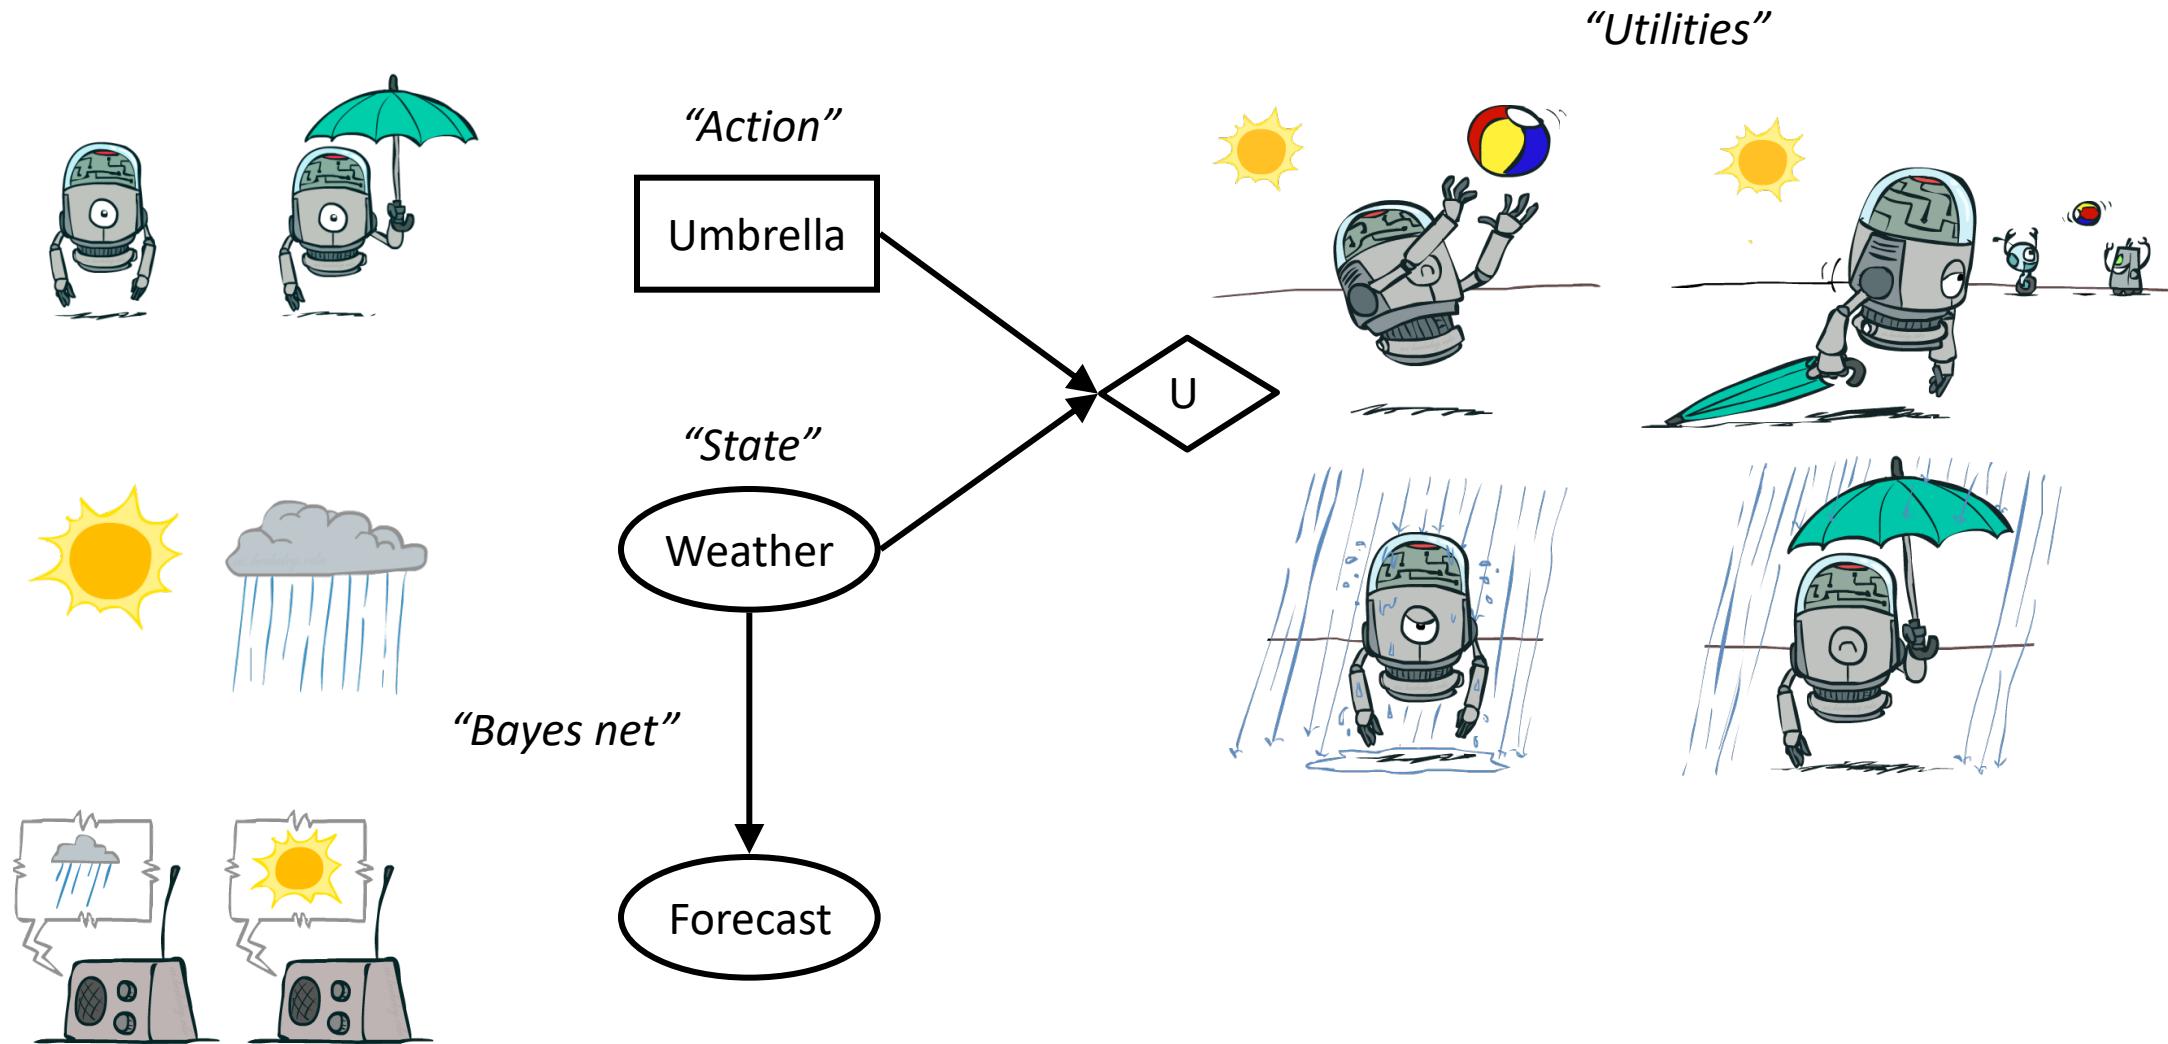
[Slides adapted from UC Berkeley CS188]



# Decision Networks

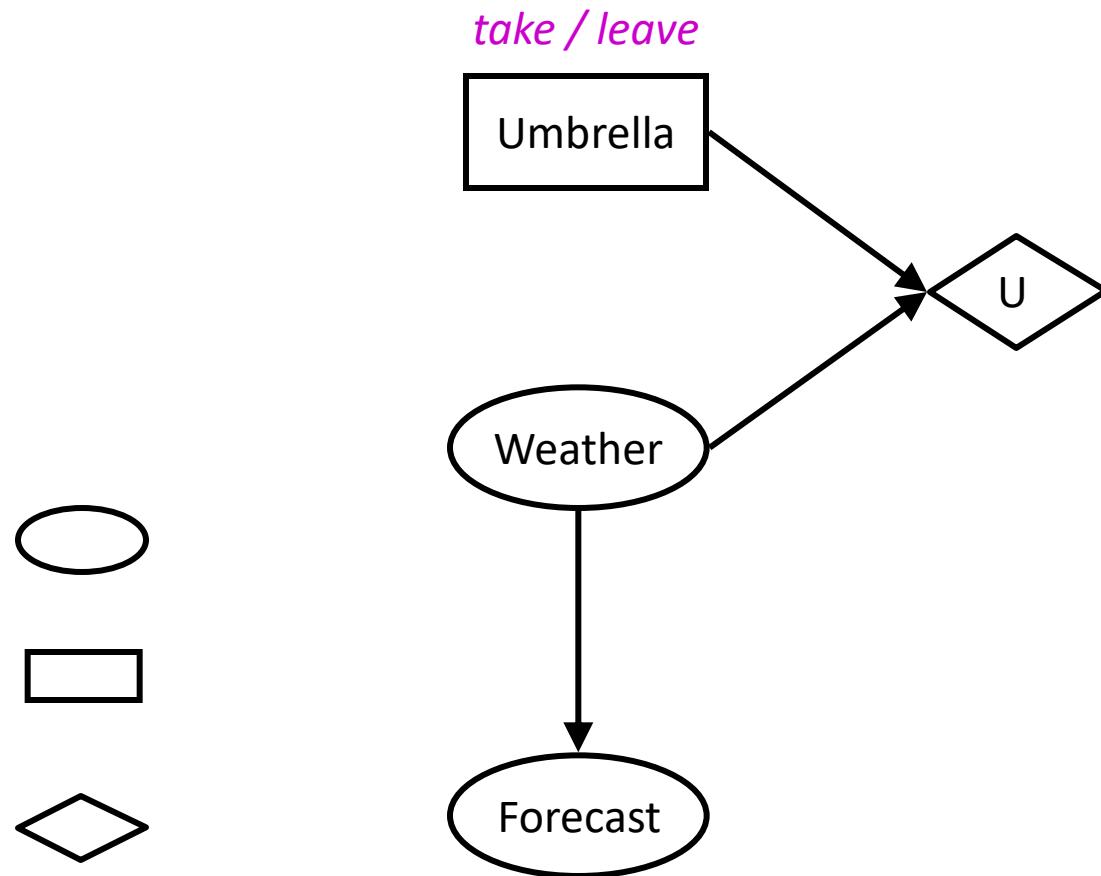


# Decision Networks



# Decision Networks

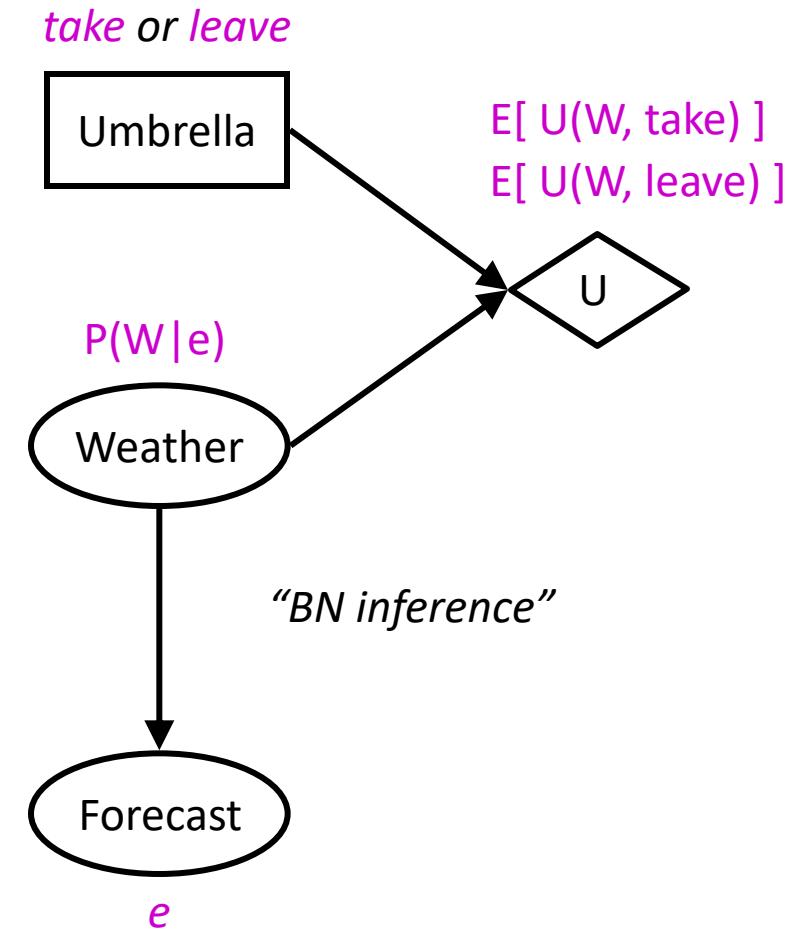
- **MEU: choose the action which maximizes the expected utility given the evidence**
- Using decision networks:
  - Bayes nets with nodes for utility and actions
  - Lets us calculate the expected utility for each action
- New node types:
  - Chance nodes (just like BNs)
  - Actions (rectangles, *cannot* have parents, act as observed evidence)
  - Utility node (diamond, depends on action and chance nodes)



# Decision Networks

- Action selection

- Instantiate all evidence
- Set action node(s) each possible way
- Calculate posterior for all parents of utility node, given the evidence
- Calculate expected utility for each action
- Choose maximizing action



# Example: Umbrella Decision Network

Umbrella = leave

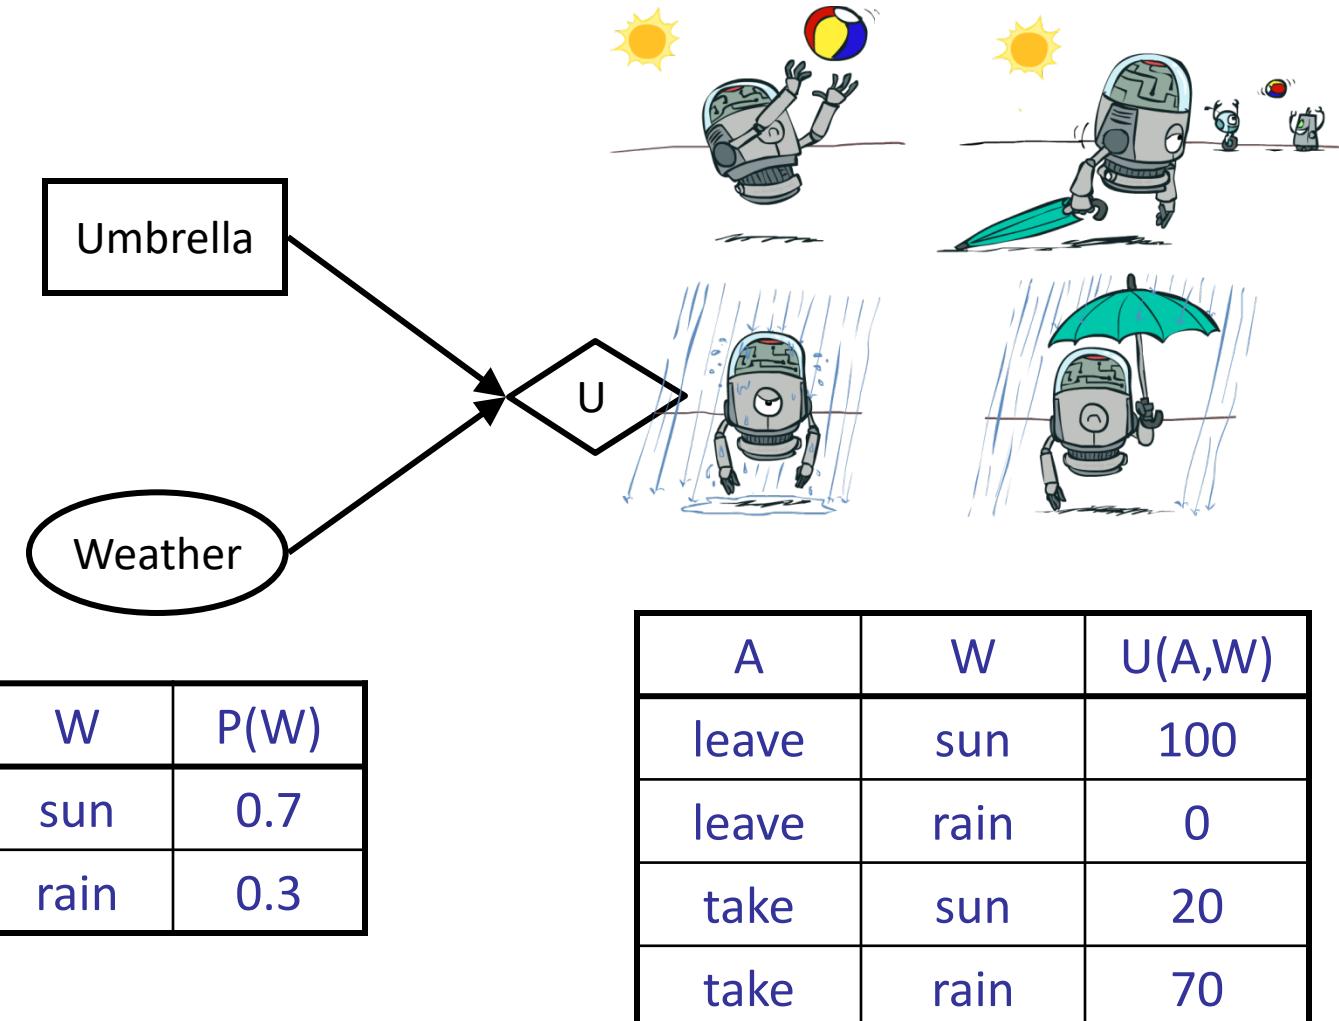
$$\begin{aligned} \text{EU}(\text{leave}) &= \sum_w P(w)U(\text{leave}, w) \\ &= 0.7 \cdot 100 + 0.3 \cdot 0 = 70 \end{aligned}$$

Umbrella = take

$$\begin{aligned} \text{EU}(\text{take}) &= \sum_w P(w)U(\text{take}, w) \\ &= 0.7 \cdot 20 + 0.3 \cdot 70 = 35 \end{aligned}$$

Optimal decision = leave

$$\text{MEU}(\emptyset) = \max_a \text{EU}(a) = 70$$



# Decision Networks: Notation

Umbrella = leave

$$\begin{aligned} \text{EU(leave)} &= \sum_w P(w)U(\text{leave}, w) \\ &= 0.7 \cdot 100 + 0.3 \cdot 0 = 70 \end{aligned}$$

Umbrella = take

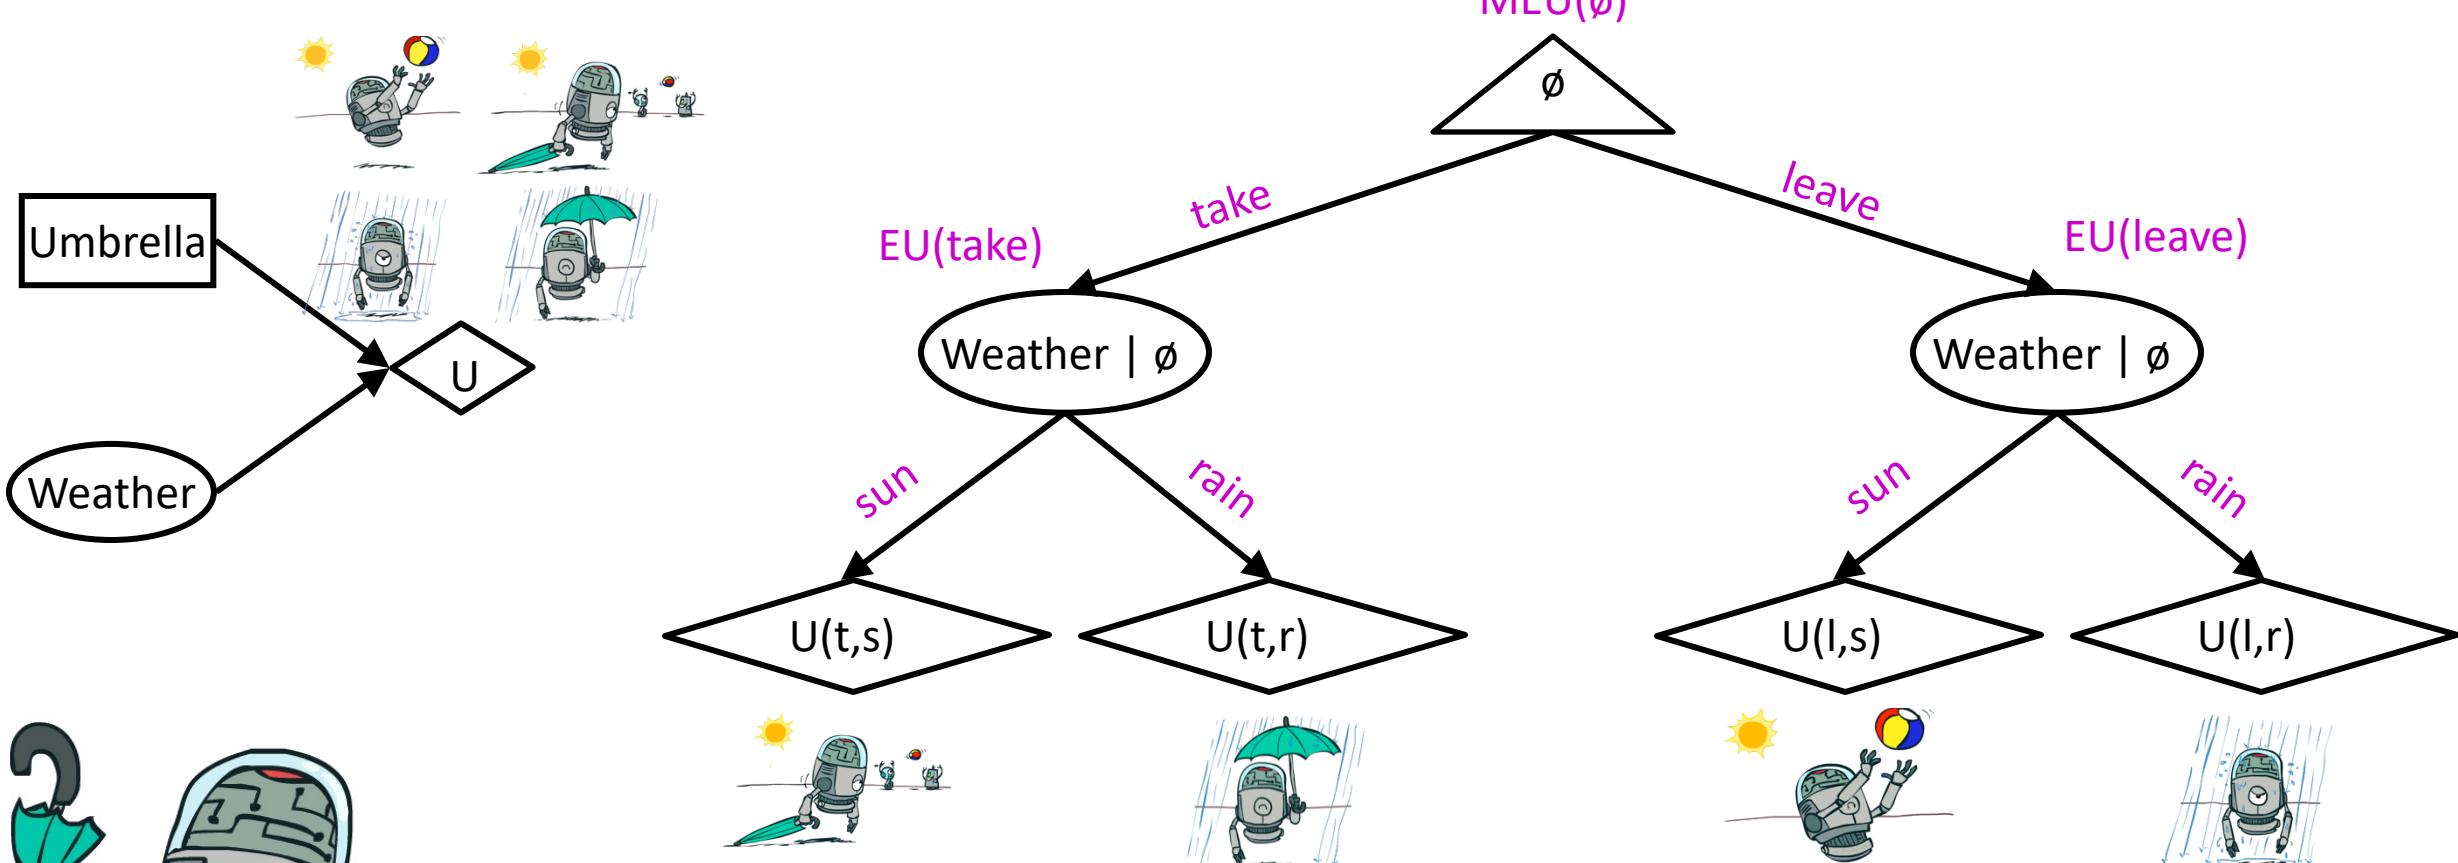
$$\begin{aligned} \text{EU(take)} &= \sum_w P(w)U(\text{take}, w) \\ &= 0.7 \cdot 20 + 0.3 \cdot 70 = 35 \end{aligned}$$

Optimal decision = leave

$$\text{MEU}(\emptyset) = \max_a \text{EU}(a) = 70$$

- EU(leave) = Expected Utility of taking action leave
  - In the parentheses, we write an action
  - Calculating EU requires taking an expectation over chance node outcomes
- MEU( $\emptyset$ ) = Maximum Expected Utility, given no information
  - In the parentheses, we write the evidence (which nodes we know)
  - Calculating MEU requires taking a maximum over several expectations (one EU per action)

# Decisions as Outcome Trees



- Almost exactly like expectimax / MDPs
- What's changed?

# Umbrella Decision Network (Evidence)

Umbrella = leave

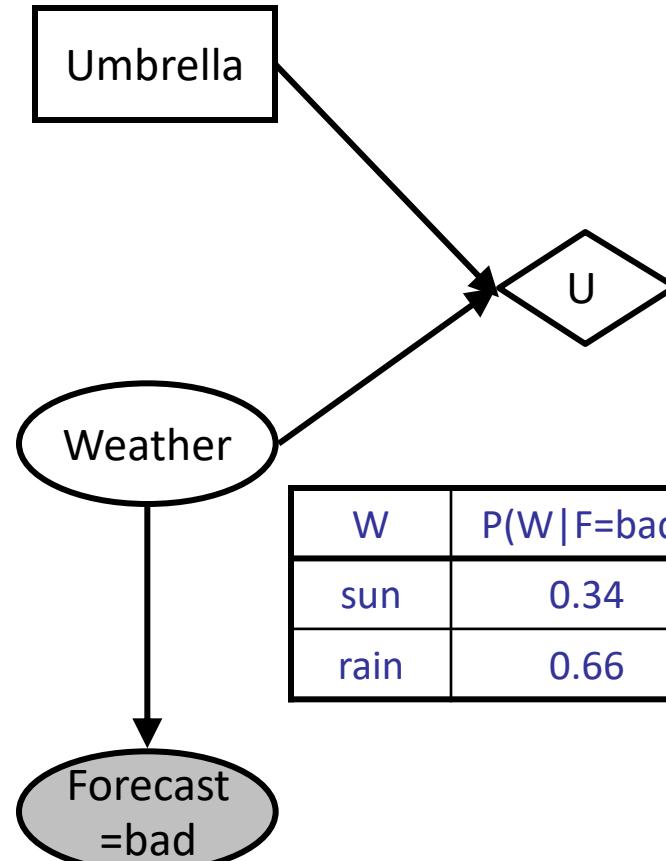
$$\begin{aligned} \text{EU}(\text{leave}|\text{bad}) &= \sum_w P(w|\text{bad})U(\text{leave}, w) \\ &= 0.34 \cdot 100 + 0.66 \cdot 0 = 34 \end{aligned}$$

Umbrella = take

$$\begin{aligned} \text{EU}(\text{take}|\text{bad}) &= \sum_w P(w|\text{bad})U(\text{take}, w) \\ &= 0.34 \cdot 20 + 0.66 \cdot 70 = 53 \end{aligned}$$

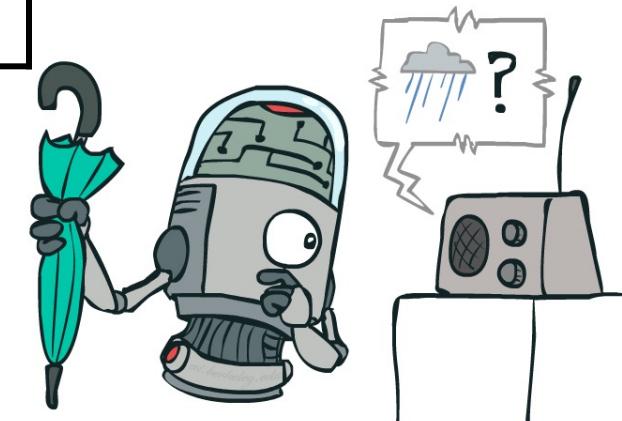
Optimal decision = take

$$\text{MEU}(F = \text{bad}) = \max_a \text{EU}(a|\text{bad}) = 53$$



W	P(W F=bad)
sun	0.34
rain	0.66

A	W	U(A,W)
leave	sun	100
leave	rain	0
take	sun	20
take	rain	70



# Decision Networks: Notation

Umbrella = leave

$$\begin{aligned} \text{EU}(\text{leave}|\text{bad}) &= \sum_w P(w|\text{bad})U(\text{leave}, w) \\ &= 0.34 \cdot 100 + 0.66 \cdot 0 = 34 \end{aligned}$$

Umbrella = take

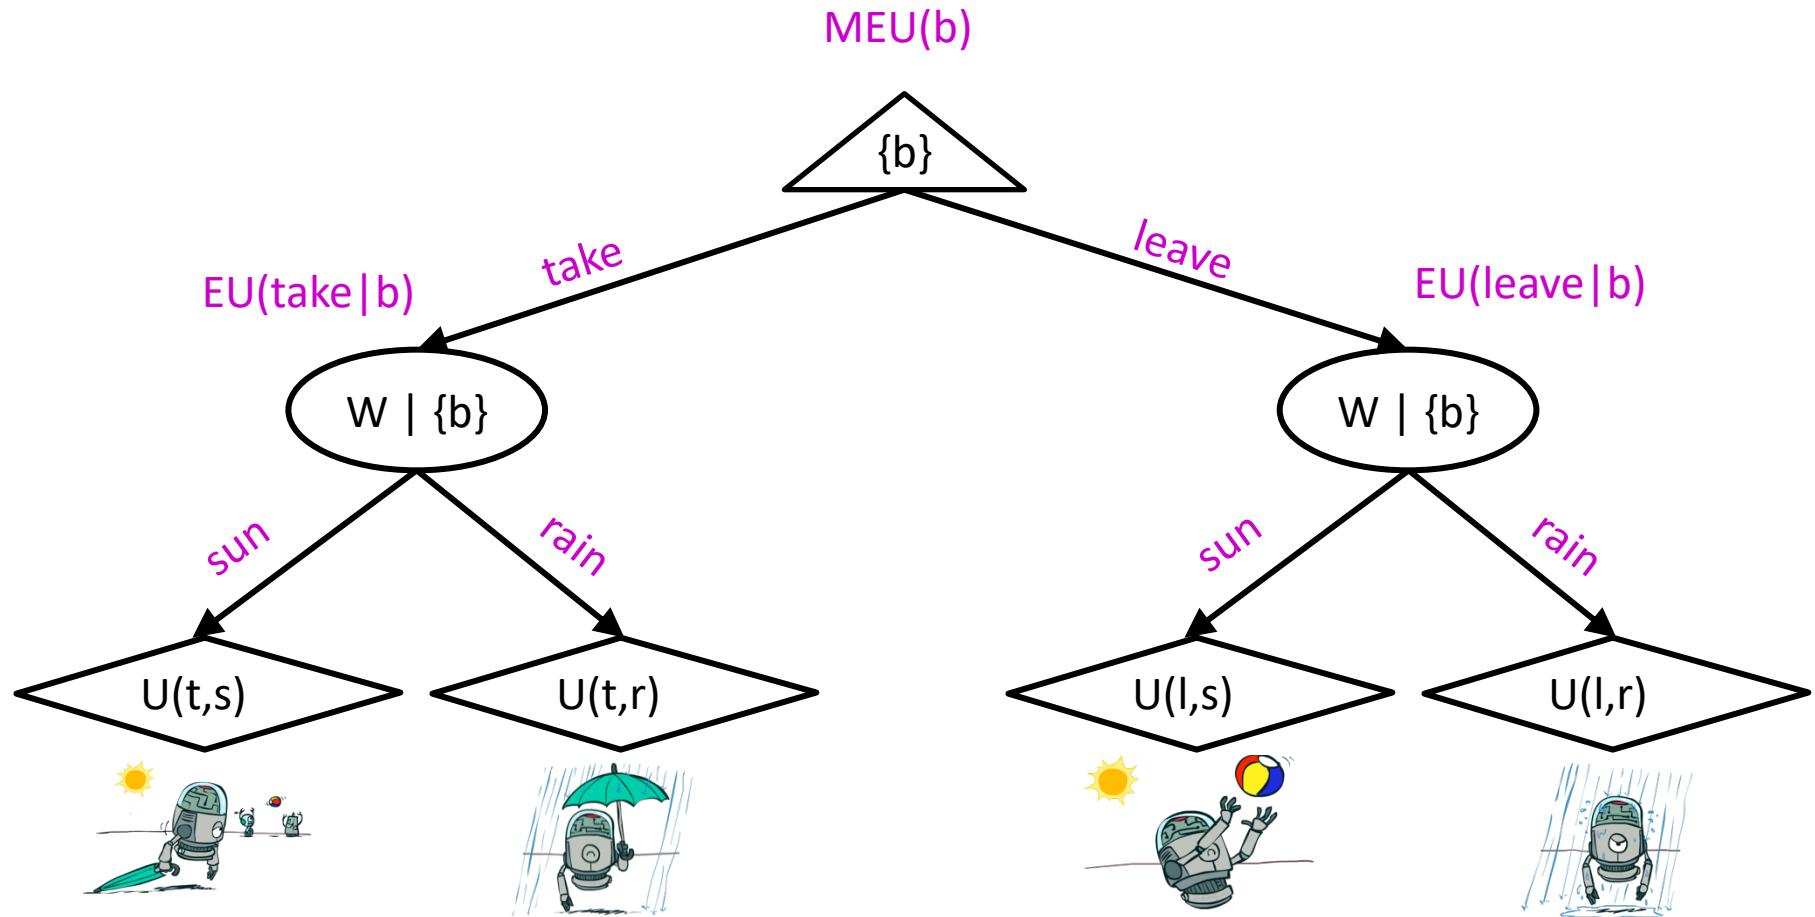
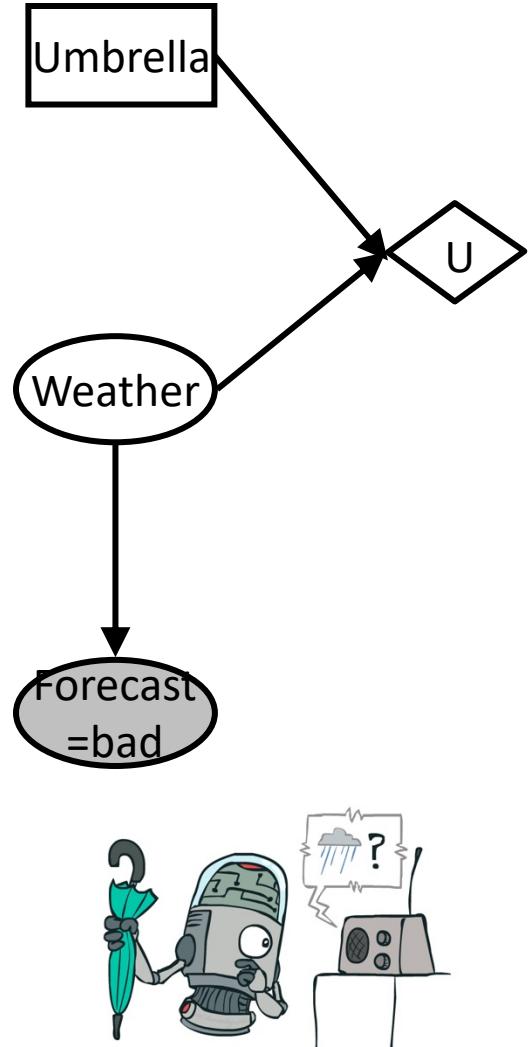
$$\begin{aligned} \text{EU}(\text{take}|\text{bad}) &= \sum_w P(w|\text{bad})U(\text{take}, w) \\ &= 0.34 \cdot 20 + 0.66 \cdot 70 = 53 \end{aligned}$$

Optimal decision = take

$$\text{MEU}(F = \text{bad}) = \max_a \text{EU}(a|\text{bad}) = 53$$

- **EU(leave|bad)** = Expected Utility of taking action leave, given you know the forecast is bad
  - Left side of conditioning bar: Action being taken
  - Right side of conditioning bar: The random variable(s) we know the value of (evidence)
- **MEU(F=bad)** = Maximum Expected Utility, given you know the forecast is bad
  - In the parentheses, we write the evidence (which nodes we know)
- Remind you of anything?

# Decisions as Outcome Trees (Evidence)

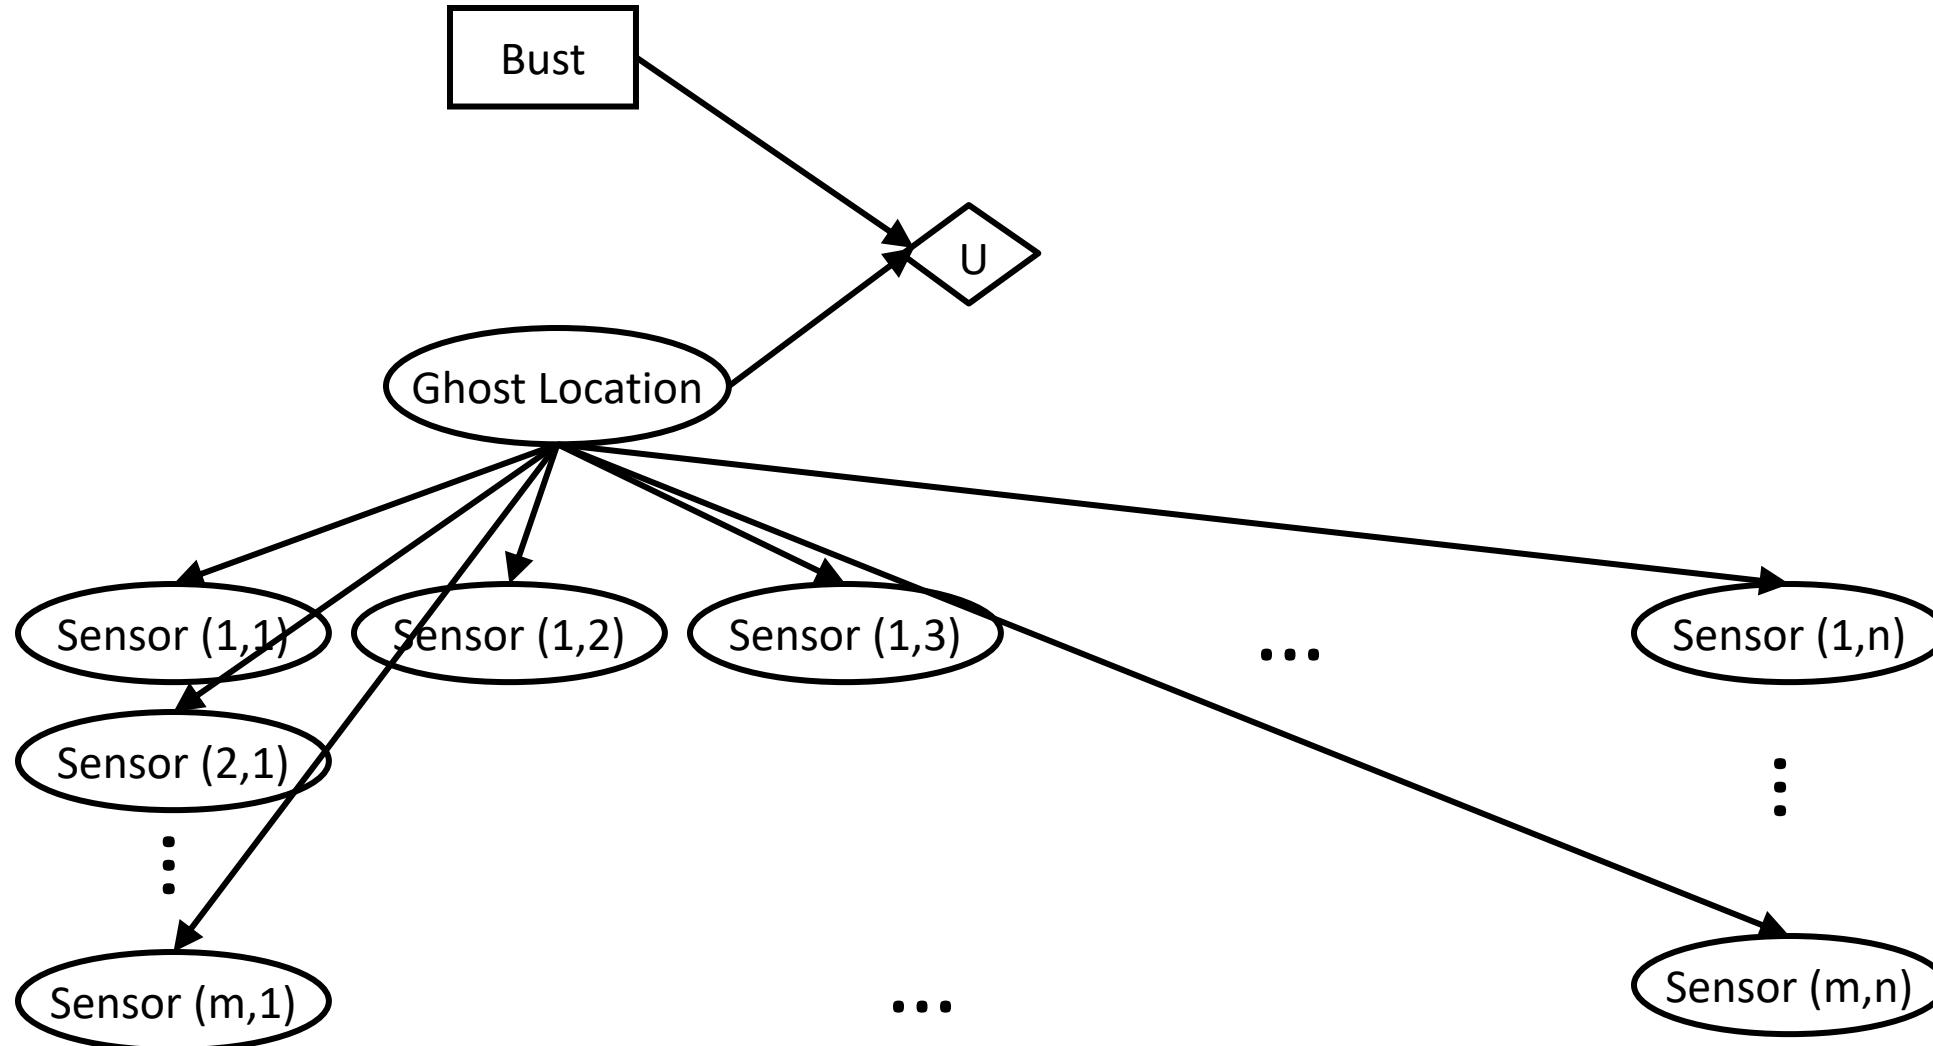


# Video of Demo Ghostbusters with Probability

---

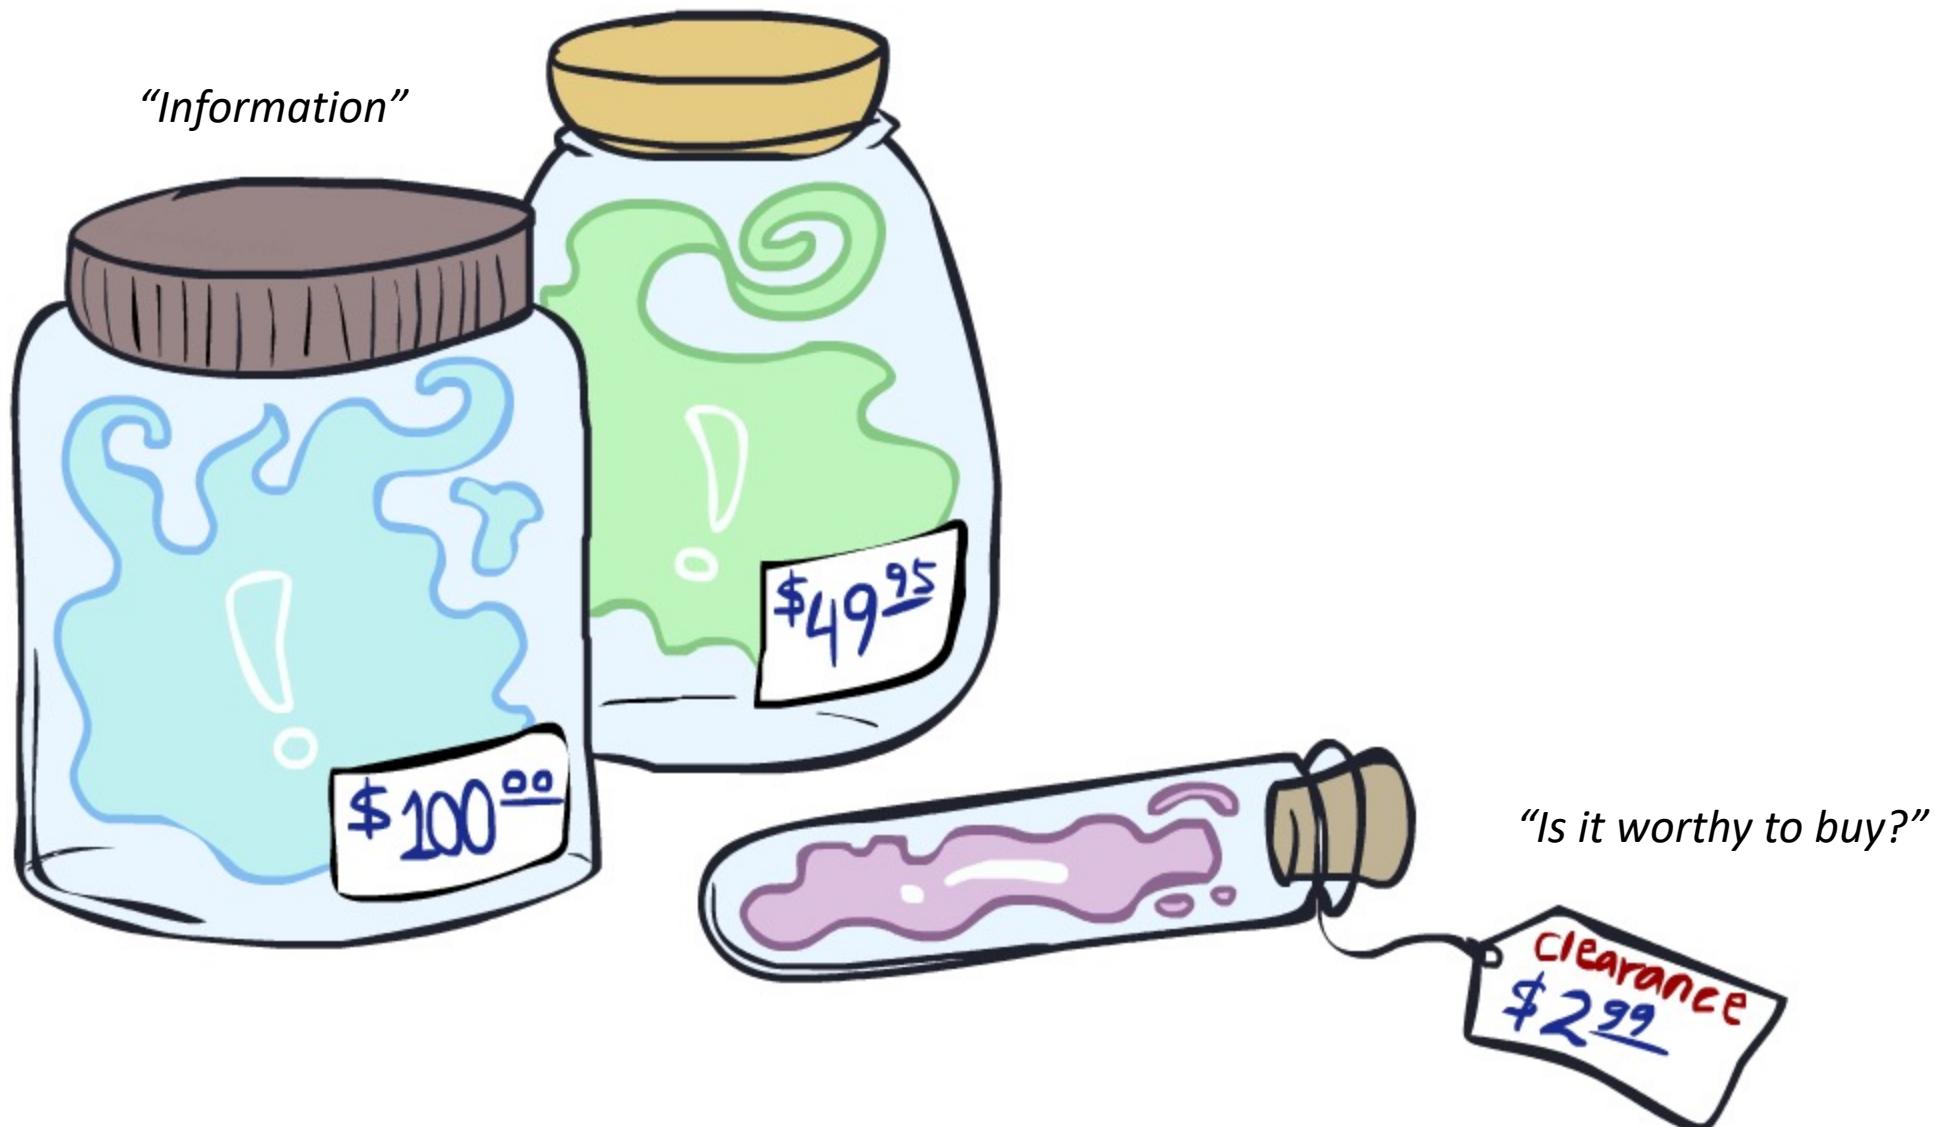


# Ghostbusters Decision Network



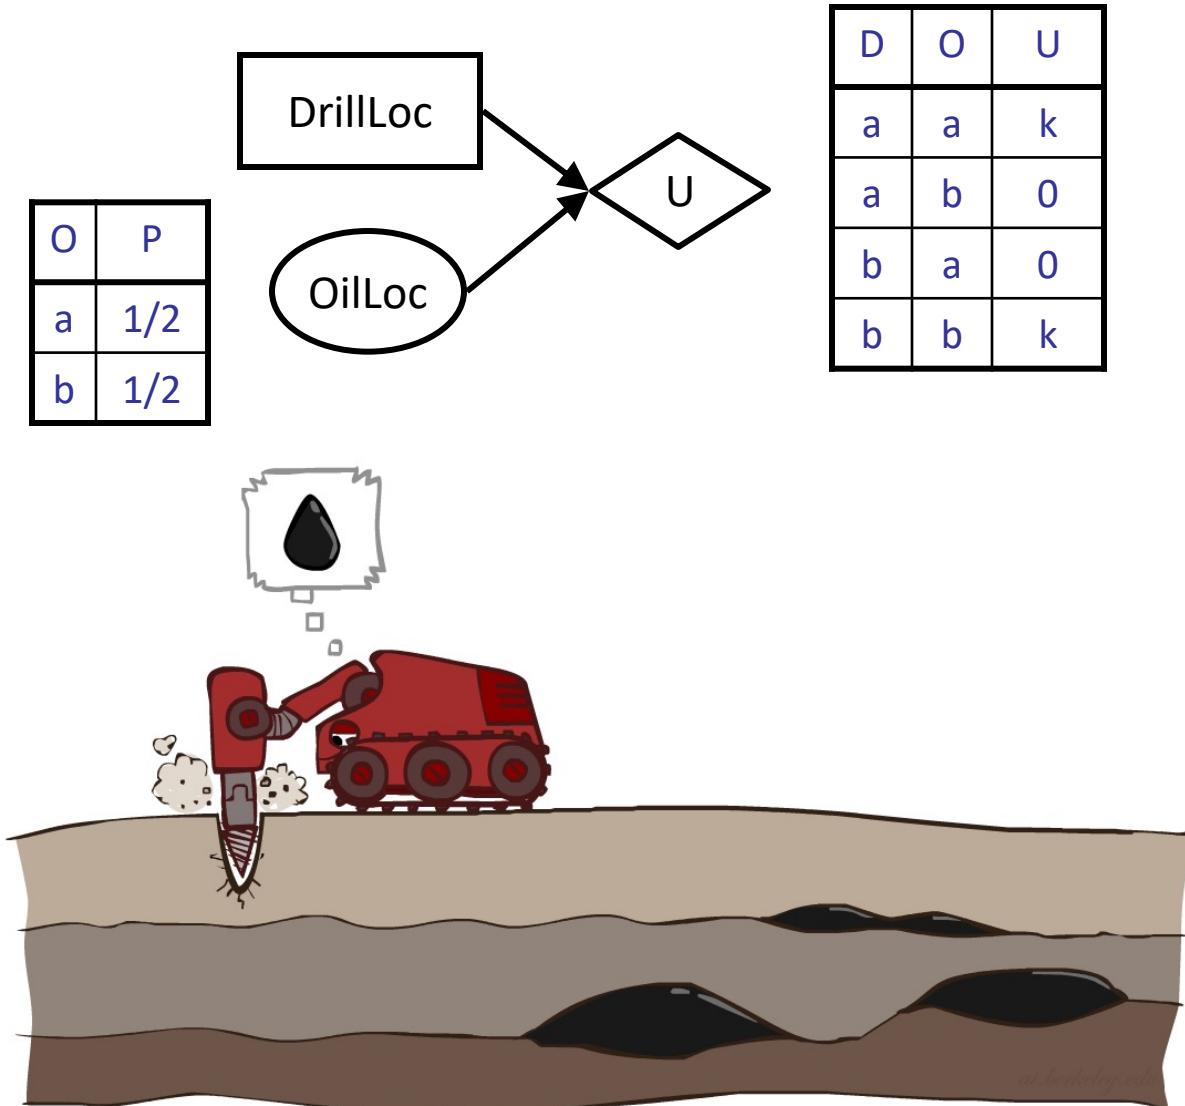
# Value of Information

---



# Value of Information

- Idea: compute value of acquiring evidence
  - Can be done directly from decision network
- Example: buying oil drilling rights
  - Two blocks A and B, exactly one has oil, worth k
  - You can drill in one location
  - Prior probabilities 0.5 each, & mutually exclusive
  - Drilling in either A or B has  $EU = k/2$ ,  $MEU = k/2$
- Question: what's the value of information of O?
  - Value of knowing which of A or B has oil from a survey
  - Value is expected gain in MEU from new info
  - Survey may say "oil in a" or "oil in b", prob 0.5 each
  - If we know OilLoc, MEU is k (either way)
  - Gain in MEU from knowing OilLoc?
  - $VPI(OilLoc) = k/2$
  - Fair price of information:  $k/2$



# VPI Example: Weather

MEU with no evidence

$$\text{MEU}(\emptyset) = \max_a \text{EU}(a) = 70$$

MEU if forecast is bad

$$\text{MEU}(F = \text{bad}) = \max_a \text{EU}(a|\text{bad}) = 53$$

MEU if forecast is good

$$\text{MEU}(F = \text{good}) = \max_a \text{EU}(a|\text{good}) = 95$$

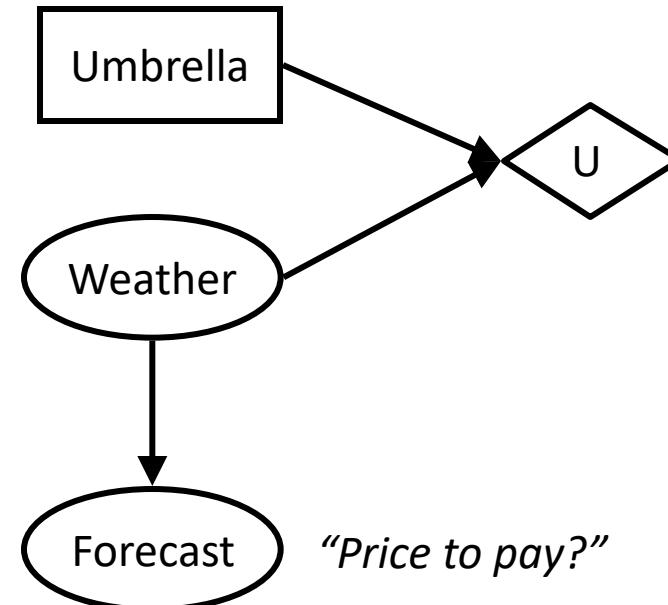
Forecast distribution

F	P(F)
good	0.59
bad	0.41

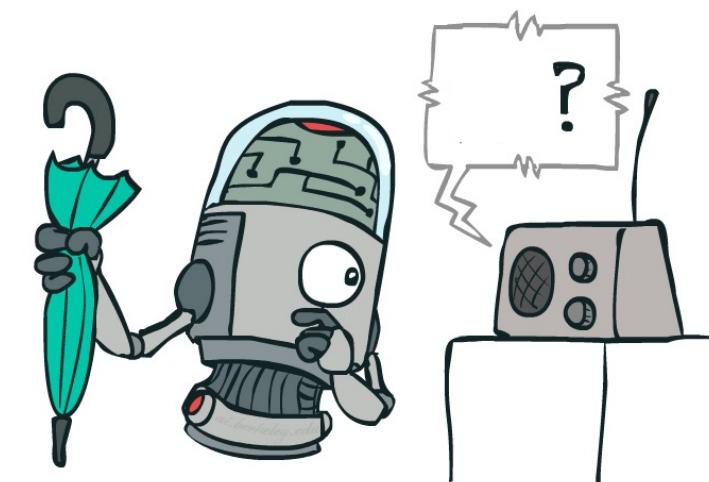


$$0.59 \cdot (95) + 0.41 \cdot (53) - 70 \\ 77.8 - 70 = 7.8$$

$$\text{VPI}(E'|e) = \left( \sum_{e'} P(e'|e) \text{MEU}(e, e') \right) - \text{MEU}(e)$$



A	W	U
leave	sun	100
leave	rain	0
take	sun	20
take	rain	70



# Value of Information (Breakdown)

- Assume we have evidence  $E=e$ . MEU if we act now:

$$\text{MEU}(e) = \max_a \sum_s P(s|e) U(s, a)$$

- Assume we see that  $E' = e'$ . MEU if we act then:

$$\text{MEU}(e, e') = \max_a \sum_s P(s|e, e') U(s, a)$$

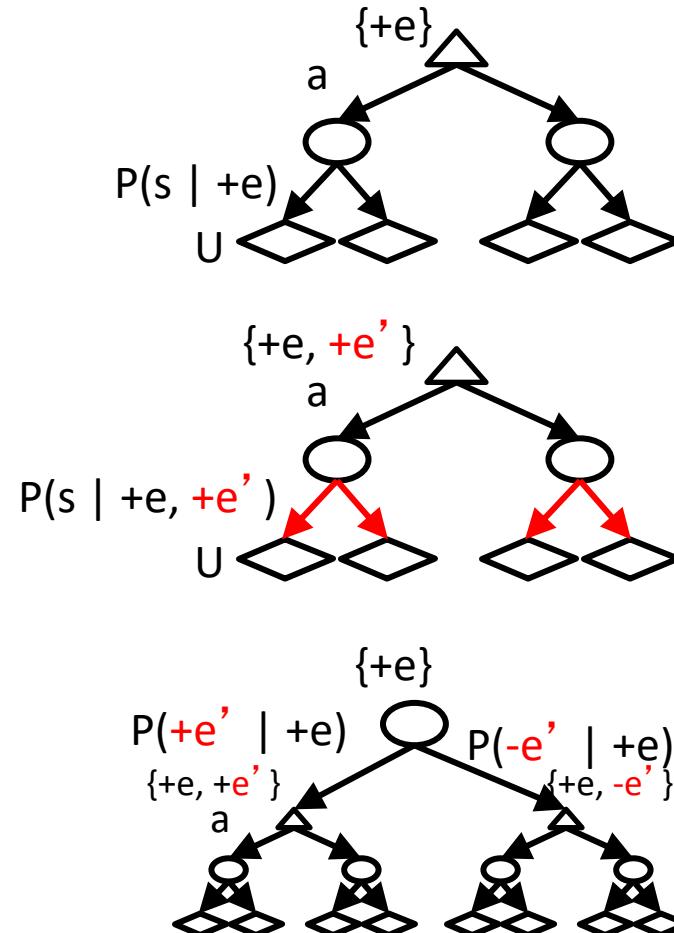
- BUT  $E'$  is a random variable whose value is *unknown*, so we don't know what  $e'$  will be

- Expected MEU if  $E'$  is revealed and then we act:

$$\text{MEU}(e, E') = \sum_{e'} P(e'|e) \text{MEU}(e, e')$$

- Value of information: how much MEU goes up by revealing  $E'$  first then acting, over acting now:

$$\text{VPI}(E'|e) = \text{MEU}(e, E') - \text{MEU}(e)$$



# VPI: Notation

- **MEU( $e$ ) = Maximum Expected Utility, given evidence  $E=e$** 
  - In the parentheses, we write the evidence (which nodes we know)
  - Calculating MEU requires taking a maximum over several expectations (one EU per action)
- **VPI( $E'|e$ ) = Expected gain in utility for knowing the value of  $E'$ , given that I already know the value of  $e$  so far**
  - Left side of conditioning bar: The random variable(s) we want to know the value of revealing
  - Right side of conditioning bar: The random variable(s) we already know the value of
  - Calculating VPI requires taking an expectation over several MEUs (one MEU per possible outcome of  $E'$ , because we don't know the value of  $E'$ )

$$\text{MEU}(e) = \max_a \sum_s P(s|e) U(s, a)$$

$$\text{MEU}(e, e') = \max_a \sum_s P(s|e, e') U(s, a)$$

$$\text{VPI}(E'|e) = \left( \sum_{e'} P(e'|e) \text{MEU}(e, e') \right) - \text{MEU}(e)$$

# VPI: Computation Workflow

---

$$\text{MEU}(e, E') \quad \text{MEU}(e, E') = \sum_{e'} P(e'|e) \text{MEU}(e, e') \quad \text{MEU}(e, e') = \max_a \text{EU}(a)$$

$$- \text{MEU}(e)$$

$$\text{MEU}(e) = \max_a \text{EU}(a)$$

$$= \text{VPI}(E'|e)$$

# VPI Properties

- Nonnegative

$$\forall E', e : \text{VPI}(E'|e) \geq 0$$



- Nonadditive

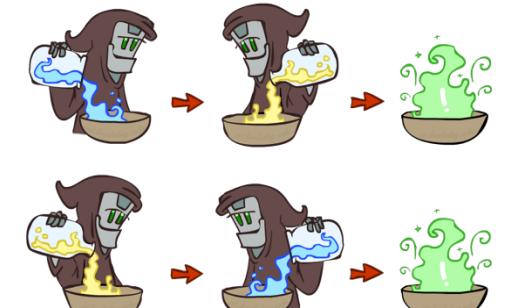
(think of observing  $E_j$  twice)

$$\text{VPI}(E_j, E_k|e) \neq \text{VPI}(E_j|e) + \text{VPI}(E_k|e)$$



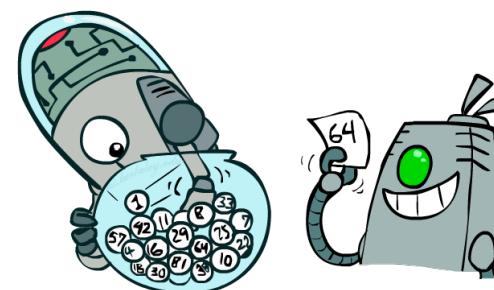
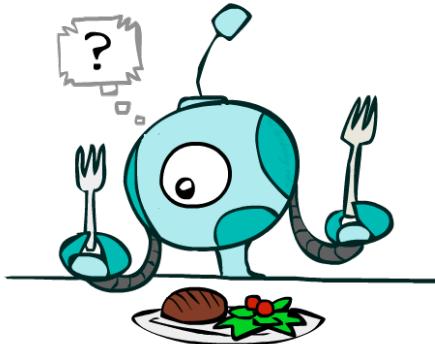
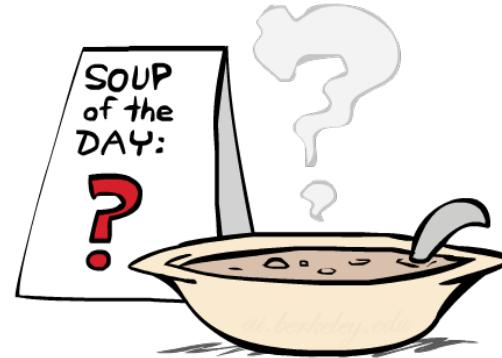
- Order-independent

$$\begin{aligned} \text{VPI}(E_j, E_k|e) &= \text{VPI}(E_j|e) + \text{VPI}(E_k|e, E_j) \\ &= \text{VPI}(E_k|e) + \text{VPI}(E_j|e, E_k) \end{aligned}$$



# Quick VPI Questions

- The soup of the day is either clam chowder or split pea, but you wouldn't order either one. What's the value of knowing which it is?
- There are two kinds of plastic forks at a picnic. One kind is slightly sturdier. What's the value of knowing which?
- You're playing the lottery. The prize will be \$0 or \$100. You can play any number between 1 and 100 (chance of winning is 1%). What is the value of knowing the winning number?



# Value of Imperfect Information?

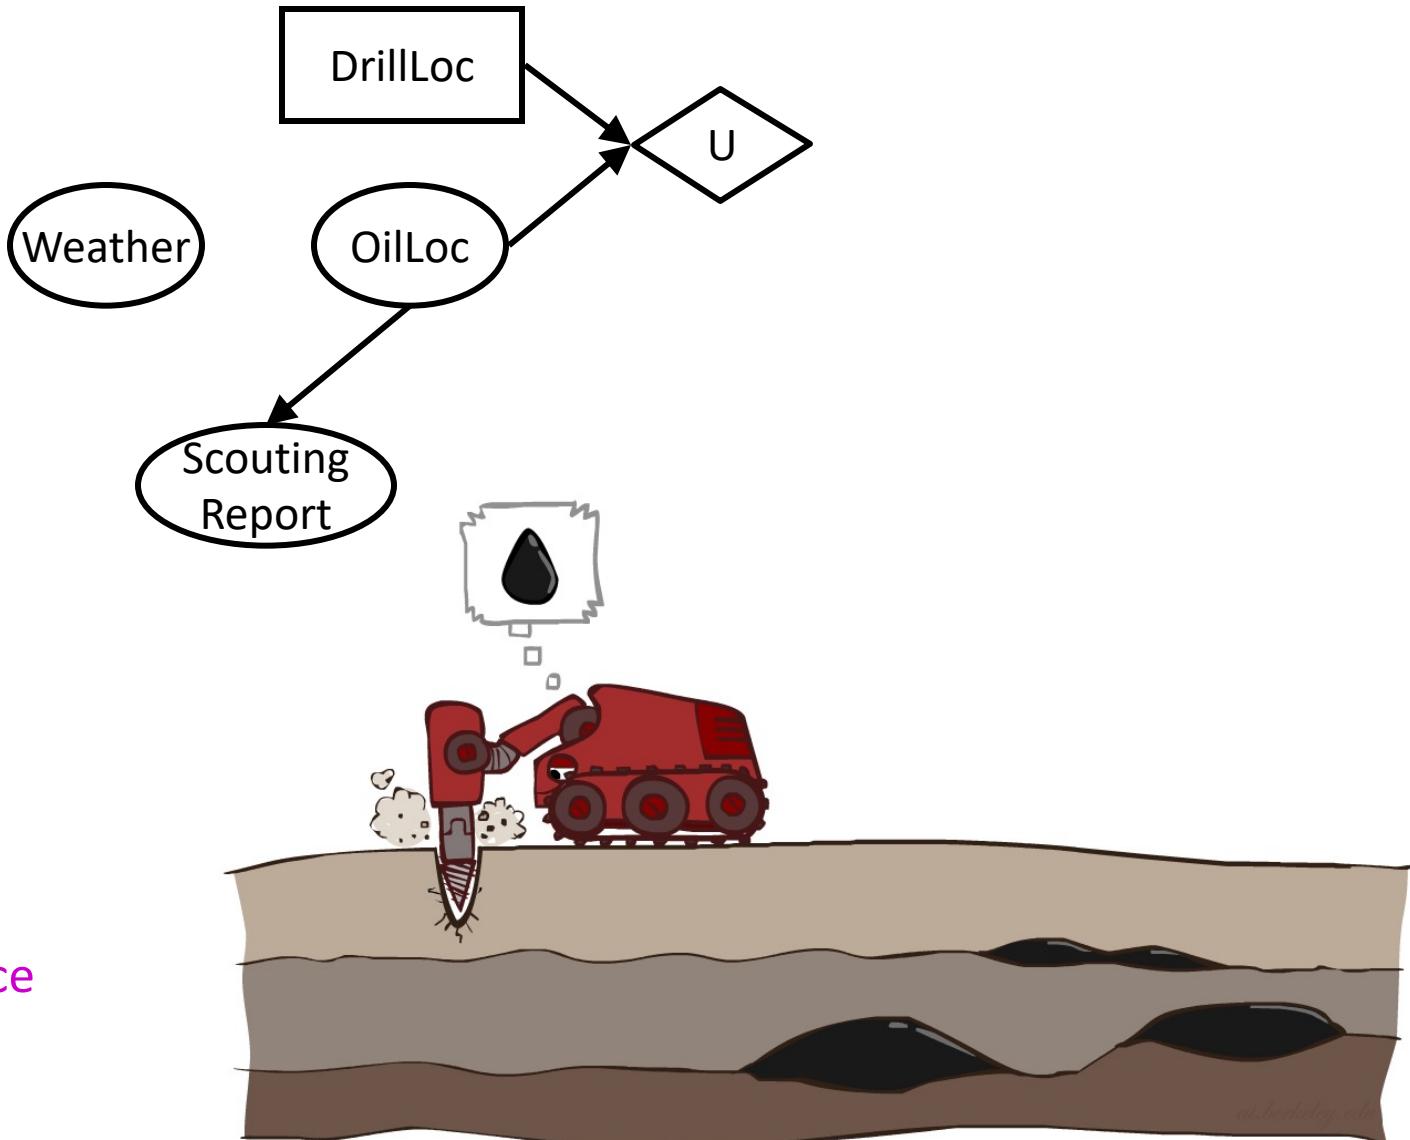
---



- No such thing (as we formulate it)
- Information corresponds to the observation of a node in the decision network
- If data is “noisy” that just means we don’t observe the original variable, but another variable which is a noisy version of the original one

# VPI Question: Oil Net

- VPI(OilLoc) ?
- VPI(ScoutingReport) ?
- VPI(Weather) ?
- VPI(Weather | ScoutingReport) ?
- Generally:  
If  $\text{Parents}(U) \perp\!\!\!\perp Z \mid \text{CurrentEvidence}$   
Then  $\text{VPI}(Z \mid \text{CurrentEvidence}) = 0$



# Next Time: Hidden Markov Models

---

# MCTS Pseudo-code (MDP)

---

**Algorithm 4.9** Monte Carlo tree search

---

```
1: function SELECTACTION( $s, d$ )           s: state; d: remaining steps of episode
2:   loop
3:     SIMULATE( $s, d, \pi_0$ )             e.g., try 1000 simulations to estimate  $Q(s, a)$ 
4:   return  $\arg \max_a Q(s, a)$ 
5: function SIMULATE( $s, d, \pi_0$ )
6:   if  $d = 0$                          Termination of a simulation (episode length exhausted, terminal state)
7:     return 0
8:   if  $s \notin T$                    If s is not in the tree (new node)
9:     for  $a \in A(s)$ 
10:       $(N(s, a), Q(s, a)) \leftarrow (N_0(s, a), Q_0(s, a))$           Initialize quantities in the new node
11:       $T = T \cup \{s\}$                 Add to the tree, then rollout
12:      return ROLLOUT( $s, d, \pi_0$ )
13:       $a \leftarrow \arg \max_{a \in A(s)} \left[ Q(s, a) + c \sqrt{\frac{\log N(s)}{N(s, a)}} \right]$     Select action branch to traverse (for both new/existing nodes)
14:       $(s', r) \sim G(s, a)$             Sample transition and reward (G: step function)
15:       $q \leftarrow r + \gamma \text{SIMULATE}(s', d - 1, \pi_0)$         Forward recursion and approximate Q-update
16:       $N(s, a) \leftarrow N(s, a) + 1$ 
17:       $Q(s, a) \leftarrow Q(s, a) + \frac{q - Q(s, a)}{N(s, a)}$         Backup: updating quantities in the node (running average)
18:      return  $q$                       Return the outcome of the current simulation
```

---

# MCTS Pseudo-code (MDP)

---

## Algorithm 4.10 Rollout evaluation

---

```
1: function ROLLOUT( $s, d, \pi_0$ )
2:   if  $d = 0$            Termination of rollout (length exhausted or terminal state)
3:     return 0
4:    $a \sim \pi_0(s)$       Sample action from rollout policy
5:    $(s', r) \sim G(s, a)$  Sample transition and reward from step function
6:   return  $r + \gamma \text{ROLLOUT}(s', d - 1, \pi_0)$  Forward recursion and back up (accumulate rewards)
```

---

# MCTS Pseudo-code (MiniMax)

```
1: function SELECT_ACTION( $s$ )           s: state of the board
2:   loop
3:     SIMULATE( $s, \pi_0, \text{player}$ )      e.g., try 1000 simulations to estimate  $U(s')$ , player = MAX
4:   return  $\arg \max_s U(s')$ 

5: function SIMULATE( $s, \pi_0, \text{player}$ )
6:   if player wins or draw                Termination of a simulation (win or draw)
7:     return 1 (MAX) or -1 (MIN) or 0 (draw)
8:   if  $s$  not in  $T$                       If  $s$  is not in the tree (new node)
9:      $U(s) = U_0(s), N(s) = N_0(s)$           Initialize quantities in the new node
10:     $T = T \cup \{s\}$ 
11:    return ROLLOUT( $s, \pi_0, \text{player}$ )    Add to the tree, then rollout
12:     $s' \leftarrow \text{random(unexplored successors)}$  if  $s$  is leaf
13:     $s' \leftarrow \arg \max_{s'} [ U(s') + c * \sqrt{\log(N(s)) / N(s')} ]$  if MAX player    Select next state to traverse (choose 1 case)
14:     $s' \leftarrow \arg \min_{s'} [ U(s') - c * \sqrt{\log(N(s)) / N(s')} ]$  if MIN player
15:     $r \leftarrow \text{SIMULATE}(s', \pi_0, -\text{player})$           Forward recursion and update utility estimate
16:     $N(s) \leftarrow N(s) + 1$ 
17:     $U(s) \leftarrow U(s) + (r - U(s)) / N(s)$           Backup: updating quantities in the node (running average)
18:  return  $r$                                 Return the outcome of the current simulation
```

# MCTS Pseudo-code (MiniMax)

---

```
1: function ROLLOUT( $s, \pi_0, \text{player}$ )
2:   if  $s$  is a terminal state
3:     return  $+1/-1/0$ 
4:    $s' \leftarrow \pi_0(s)$ 
5:   return ROLLOUT( $s', \pi_0, -\text{player}$ )
```

*Termination of rollout*

*Sample plays from rollout policy*

*Forward recursion and back up (accumulate rewards)*