


CS 3317: Artificial Intelligence

Particle Filters




Instructor: Panpan Cai

[Slides adapted from UC Berkeley CS188]




Hidden Markov Models

- Markov chains not so useful for most agents
 - Real-problems are often partially observable
 - Use observations to update your beliefs
- Hidden Markov models (HMMs)
 - Underlying Markov chain over states X_i
 - You observe evidences at each time step





Example: Ghostbusters HMM

- $P(X_0)$ = uniform
- $P(X' | X)$ = usually move clockwise, but sometimes move in a random direction or stay in place
- $P(R_{ij} | X')$ = same sensor model as before:
red means close, green means far away.



1/9	1/9	1/9
1/9	1/9	1/9
1/9	1/9	1/9

$P(X_0)$



1/6	1/6	1/2
0	1/6	0
0	0	0

$P(X' | X=<1,2>)$

Filtering / Belief Tracking

- Filtering is the task of tracking the *belief state*

$$B_t(X) = P(X_t \mid e_1, \dots, e_t)$$


over time

- We start with $B_0(X)$ in an initial setting, usually *uniform*
- As time passes, or when we get observations, we update $B(X)$

Ghostbusters Basic Dynamics



Ghostbusters – Circular Dynamics -- HMM



Ghostbusters Circular Dynamics (No-observation)



Ghostbusters Whirlpool Dynamics



Recursive Filtering

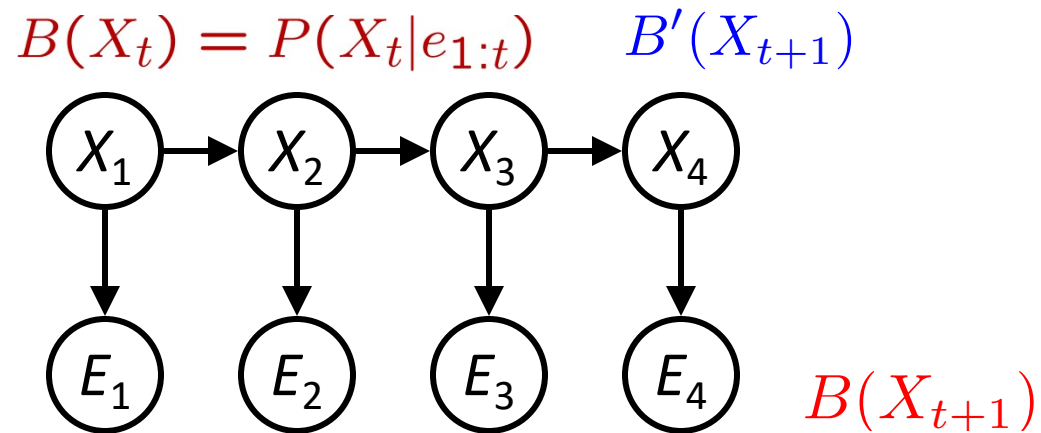
- We are given evidence at each time till now and want to know

$$B_t(X) = P(X_t | e_{1:t})$$

- Idea of recursive filtering:
 - Start with $P(X_0)$
 - Derive B_t in terms of B_{t-1} given e_t
 - Equivalently, derive B_{t+1} in terms of B_t given e_{t+1}

Two Steps


1. **Passage of time**: predict possible transitions
2. **Observation**: update by incorporating the observation



Step 1: Passage of Time

- Assume we have current belief $P(X \mid \text{evidence to date})$

$$B(X_t) = P(X_t | e_{1:t})$$



- After one time step passes:

$$\begin{aligned} P(X_{t+1} | e_{1:t}) &= \sum_{x_t} P(X_{t+1}, x_t | e_{1:t}) \\ &= \sum_{x_t} P(X_{t+1} | x_t, e_{1:t}) P(x_t | e_{1:t}) \\ &= \sum_{x_t} P(X_{t+1} | x_t) P(x_t | e_{1:t}) \end{aligned}$$

- Or compactly:

$$B'(X_{t+1}) = \sum_{x_t} P(X' | x_t) B(x_t)$$

- Basic idea: beliefs get “pushed” through the transitions
 - With the “B” notation, we have to be careful about what time step t the belief is about, and what evidence it includes


Example: Passage of Time

- As time passes, uncertainty “accumulates”

(Transition model: ghosts usually go clockwise)


<0.01	<0.01	<0.01	<0.01	<0.01	<0.01
<0.01	<0.01	<0.01	<0.01	<0.01	<0.01
<0.01	<0.01	1.00	<0.01	<0.01	<0.01
<0.01	<0.01	<0.01	<0.01	<0.01	<0.01

T = 0



<0.01	<0.01	<0.01	<0.01	<0.01	<0.01
<0.01	<0.01	0.06	<0.01	<0.01	<0.01
<0.01	0.76	0.06	0.06	<0.01	<0.01
<0.01	<0.01	0.06	<0.01	<0.01	<0.01

T = 1



0.05	0.01	0.05	<0.01	<0.01	<0.01
0.02	0.14	0.11	0.35	<0.01	<0.01
0.07	0.03	0.05	<0.01	0.03	<0.01
0.03	0.03	<0.01	<0.01	<0.01	<0.01

T = 4



Step 2: Observation

- Assume we have the belief after passage of time:


$$B'(X_{t+1}) = P(X_{t+1}|e_{1:t})$$

- After evidence comes in:

$$\begin{aligned} P(X_{t+1}|e_{1:t+1}) &= P(X_{t+1}, e_{t+1}|e_{1:t}) / P(e_{t+1}|e_{1:t}) \\ &\propto_{X_{t+1}} P(X_{t+1}, e_{t+1}|e_{1:t}) \\ &= P(e_{t+1}|e_{1:t}, X_{t+1}) P(X_{t+1}|e_{1:t}) \\ &= P(e_{t+1}|X_{t+1}) P(X_{t+1}|e_{1:t}) \end{aligned}$$

- Or, compactly:

$$B(X_{t+1}) \propto_{X_{t+1}} P(e_{t+1}|X_{t+1}) B'(X_{t+1})$$



- Basic idea: beliefs “reweighted” by likelihood of evidence
- Unlike passage of time, we have to renormalize

Example: Observation

- As we get observations, beliefs get reweighted, uncertainty “decreases”


0.05	0.01	0.05	<0.01	<0.01	<0.01
0.02	0.14	0.11	0.35	<0.01	<0.01
0.07	0.03	0.05	<0.01	0.03	<0.01
0.03	0.03	<0.01	<0.01	<0.01	<0.01

Before observation

<0.01	<0.01	<0.01	<0.01	0.02	<0.01
<0.01	<0.01	<0.01	0.83	0.02	<0.01
<0.01	<0.01	0.11	<0.01	<0.01	<0.01
<0.01	<0.01	<0.01	<0.01	<0.01	<0.01

After observation

$$B(X) \propto P(e|X)B'(X)$$



Online Belief Updates



- Every time step, we start with current $P(X \mid \text{evidence})$
- We update for time:

$$P(x_t | e_{1:t-1}) = \sum_{x_{t-1}} P(x_{t-1} | e_{1:t-1}) \cdot P(x_t | x_{t-1})$$

- We update for evidence:

$$P(x_t | e_{1:t}) \propto_X P(x_t | e_{1:t-1}) \cdot P(e_t | x_t)$$

- The **FORWARD** algorithm does both at once (and doesn't normalize)



Forward Algorithm


- Transition matrix T , observation matrix O_t
 - Observation matrix contains likelihoods for E_t along its diagonal
 - E.g., for $U_1 = \text{true}$, $O_1 = \begin{pmatrix} 0.2 & 0 \\ 0 & 0.9 \end{pmatrix}$
- Forward algorithm becomes:
 - $B_{t+1} = \alpha O_{t+1} T^T B_t$
 - easy to implement in Python or MATLAB
 - *lazy* normalization

X_{t-1}	$P(X_t X_{t-1})$	
	sun	rain
sun	0.9	0.1
rain	0.3	0.7

W_t	$P(U_t W_t)$	
	true	false
sun	0.2	0.8
rain	0.9	0.1

Forward Algorithm Complexity


- $P(X_{t+1} | e_{1:t+1}) = \alpha P(e_{t+1} | X_{t+1}) \sum_{x_t} P(X_{t+1} | x_t) P(x_t | e_{1:t})$



- $B_{t+1} = \text{FORWARD}(B_t, e_{t+1})$

- Cost per time step: $O(|X|^2)$ where $|X|$ is the number of states.
- $O(|X|^2)$ is infeasible for models with large state spaces. (T_T)
- *Approximate* filtering algorithms !!! (^_^)

Particle Filtering














Particle Filtering

- Filtering: approximate solution
- Sometimes $|X|$ is too big to use exact inference
 - $|X|$ may be too big to even store $B(X)$
 - E.g. X is continuous
- Solution: approximate inference
 - Track *samples* of X , not all values
 - Samples are called *particles*
 - Time per step is linear in the *number of samples*
 - But: number needed may be large
 - In memory: list of particles, not states
- This is how robot localization works in practice
- **Particle** is just more intuitive name for sample

0.0	0.1	0.0
0.0	0.0	0.2
0.0	0.2	0.5

“Exact belief”




		
		 
	 	     

“Particle belief”

Representation: Particles

- Our representation of $P(X)$ is now a list of N particles (samples)
 - Generally, $N \ll |X|$
- $P(x)$ approximated by number of particles with value x
 - So, many x may have $P(x) = 0$!
 - More particles, more accuracy
- For now, all particles have a weight of 1

"Particle belief"



Particles:

(3,3)
(2,3)
(3,3)
(3,2)
(3,3)
(3,2)
(1,2)
(3,3)
(3,3)
(2,3)

Particle Filtering: Predict


- Each particle is moved by *sampling* its next position from the transition model

$$x' = \text{sample}(P(X'|x))$$

- This is like *prior sampling* – samples' frequencies reflect the transition probabilities
 - Ghostbuster example: most samples move clockwise, but some move in another direction or stay in place
- Predict approximates the passage of time
 - If enough samples, close to exact values (consistent)


Particles:

(3,3)
(2,3)
(3,3)
(3,2)
(3,3)
(3,2)
(1,2)
(3,3)
(3,3)
(2,3)



Particles:

(3,2)
(2,3)
(3,2)
(3,1)
(3,3)
(3,2)
(1,3)
(2,3)
(3,2)
(2,2)



Particle Filtering: Update

- Slightly trickier:

- Don't sample observation, *fix* it
- Similar to *likelihood weighting*, weight samples based on the probability of evidence


$$w(x) = P(e|x)$$

$$B(X) \propto P(e|X)B'(X)$$

- The weights don't sum to **N**, since all have been downweighted
 - In fact, they now sum to an approximation of **N*P(e)**


Particles:

(3,2)
(2,3)
(3,2)
(3,1)
(3,3)
(3,2)
(1,3)
(2,3)
(3,2)
(2,2)



Particles:

(3,2) w=.9
(2,3) w=.2
(3,2) w=.9
(3,1) w=.4
(3,3) w=.4
(3,2) w=.9
(1,3) w=.1
(2,3) w=.2
(3,2) w=.9
(2,2) w=.4



Particle Filtering: Resample



- Rather than tracking weighted samples, we *resample*
- We sample N times from the *weighted sample distribution*
- Weights of particles become 1 again
- Finishes the update for this time step, continue with the next one

Particles:

(3,2) $w=.9$
(2,3) $w=.2$
(3,2) $w=.9$
(3,1) $w=.4$
(3,3) $w=.4$
(3,2) $w=.9$
(1,3) $w=.1$
(2,3) $w=.2$
(3,2) $w=.9$
(2,2) $w=.4$


(New) Particles:

(3,2)
(2,2)
(3,2)
(2,3)
(3,3)
(3,2)
(1,3)
(2,3)
(3,2)
(3,2)



Putting Together: Particle Filtering

- Particles: track *samples* of states rather than an explicit distribution



Video of Demo – Moderate Number of Particles




Video of Demo – One Particle



Video of Demo – Huge Number of Particles

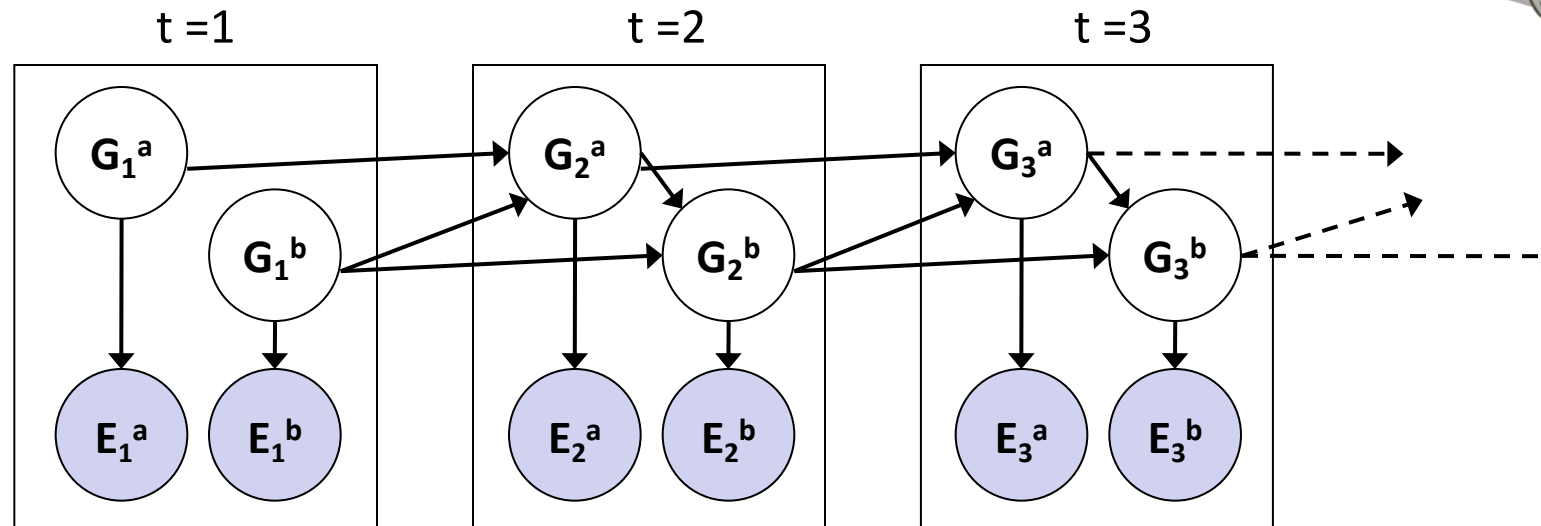



Dynamic Bayes Nets




Dynamic Bayes Nets (DBNs)

- We want to track *multiple* variables over time, using *multiple* sources of evidence
- Idea: Repeat a fixed Bayes net structure at each time
- Variables at time t can condition on those at $t-1$



DBNs and HMMs


- Every HMM is a single-variable DBN
- Every discrete DBN is an HMM
 - HMM state is Cartesian product of DBN state variables



- Sparse dependencies => exponentially fewer parameters in DBN
 - E.g., 20 Boolean state variables, 3 parents each;
DBN has $20 \times 2^3 = 160$ parameters, HMM has $2^{20} \times 2^{20} = \sim 10^{12}$ parameters

Exact Inference in DBNs

- Variable elimination applies to dynamic Bayes nets
- Offline: “unroll” the network for T time steps, then eliminate variables to find $P(X_T | e_{1:T})$





- Online: eliminate all variables from the previous time step; store factors for current time only
- Problem: largest factor contains all variables for current time (plus a few more)

DBN Particle Filters

- A particle is a *complete* sample for a time step
- **Initialize:** Generate prior samples for the $t=0$ Bayes net
 - Example particle: $\mathbf{G}_0^a = (3,3)$ $\mathbf{G}_0^b = (5,3)$
- **Predict:** Sample a successor for each particle
 - Example successor: $\mathbf{G}_1^a = (2,3)$ $\mathbf{G}_1^b = (6,3)$
- **Update:** Weight each *entire* sample by the likelihood of the evidence conditioned on the sample
 - Likelihood: $P(\mathbf{E}_1^a | \mathbf{G}_1^a) * P(\mathbf{E}_1^b | \mathbf{G}_1^b)$
- **Resample:** Select prior samples (tuples of values) in proportion to their weights

Robot Localization

- In robot localization:
 - We know the map, but not the robot's position
 - Observations may be vectors of range finder readings
 - State space and readings are typically continuous, so we cannot represent an exact belief
 - Particle filtering is a main technique




Particle Filter Localization (Sonar)




Robot Mapping

- SLAM: Simultaneous Localization And Mapping
 - Robot does not know map or location
 - State $x_t^{(i)}$ consists of position+orientation, map!
 - (Each map usually inferred exactly given sampled position+orientation sequence: RBPF)



DP-SLAM, Ron Parr



Particle Filter SLAM – Video 2

